

KNEO AGENT CLIENT

API Reference

Typed Python SDK + adapter toolkit for the Kneo Agent Platform /v1 .

Single-document reference for every public class, dataclass, exception, and helper that `kneo-client` exposes. Sibling to the `kneo_agent` SDK reference; same dark-themed layout, same hand-maintained discipline.

VERSION

v0.7.0

PINNED TO KNEO_SERV

v0.9.0 (/v1 only)

PYTHON

≥ 3.12

The user-facing prose ships separately as the Kneo Agent Client User Guide PDF. Source markdown lives under `docs/user/` on the repo. Architecture lives in `docs/dev/architecture.md` and the design handout PDF.

© 2026 Kneron, Inc. and kneo-client contributors · MIT

Table of Contents

Read the front matter first; then jump to any class. The sidebar on screen mirrors this layout; in the printed PDF, navigation is page-numbered through the per-page footer.

FRONT MATTER

What's new in v0.7.0

Introduction

TOP-LEVEL

KneoClient

CORE

Transport

SyncTransport

Profile & profiles

ApiKeyAuth & AuthScheme

RetryPolicy

Idempotency helpers

List-method results

Request-ID helpers

Logging helpers

Exceptions

PLATFORM ADAPTER

PlatformClient

HealthClient

RunsClient

HumanTasksClient

AuditClient

CredentialsClient

PoliciesClient

AGENT ADAPTER

AgentClient

SkillsClient

SpecsClient

ValidateThenCompileResult

DryRunResult

REFERENCE

Compatibility matrix

Guides

API Reference Manual

Every public class and helper, organized for both screen reading and printing. Adapted from the `kneo_agent` SDK reference's dark-themed layout; rendered to PDF via the same Puppeteer pipeline.

v0.7.0

Python ≥ 3.12

kneo_serv /v1 · 0.9.x

What's new in v0.7.0

A `kneo_serv v0.9.0` spec-pin uptake (two server minors in one cut — the largest `/v1` delta since 0.6.0) plus the resolutions of a workflow-based deep audit of the package at 0.6.0. Zero breaking changes; four narrated behavior corrections, all toward server truth.

Headline pieces:

- **Spec pin refreshed v0.7.0 → kneo_serv v0.9.0.** New `GET /v1/skills` catalog, the `RunCreateRequest.skills` overlay, the shared `environment / overlays / overrides` envelope on explain, a documented human-tasks `status` filter (retiring the contract opt-out carried since 0.3.0), `window` pagination metadata, and a structured 413 across the board.
- **agent.skills — the skill catalog.** `SkillsClient.list()` wraps the read-only catalog (`kneo_serv ≥ 0.8.0`); catalog names are the valid targets for a per-request `SkillsOverlay` on `runs.create`.
- **Page.window + honest has_more.** `total` is now the true store count and can exceed the paging window; `has_more` reports *fetchability*, so pagination loops stop at the window edge.
- **Code-aware 409s (behavior).** `run_state_conflict / resource_locked` raise `KneoConflictError` (which gains `.retry_after`); the transient `idempotency_key_in_progress` is auto-retried; only `idempotency_key_conflict` raises `KneoIdempotencyMismatchError`.
- **Typed 400 / 413 / 422.** `KneoBadRequestError`, `KneoPayloadTooLarge` (carries `max_body_bytes`), `KneoValidationError` (FastAPI list-shaped 422 detail summarized into the message).
- **Lifecycle truth.** `DEFAULT_TERMINAL_STATUSES` gains `timed_out / expired` (behavior); the phantom `paused_human_review` status is gone from the docs (the real pause status is `blocked`); `resume_first_pending()` reads the real `id` key and claims earliest-deadline-first.
- **Retry hardening.** `Retry-After` hints above `max_delay` fail fast with the hint on the typed error (behavior); transport-level errors retry only when replay-safe; auto-managed headers overwrite caller variants case-insensitively.

See the [0.7.0 release notes](#) for the full summary including the four behavior-correction upgrade notes. The [0.6.0 release notes](#) cover the prior cut's audit-pagination revival and `KneoError.code`.

Introduction

`kneo-client` is the shared client layer behind **Kneo Agent Dashboard** (operations) and **Kneo Agent Studio** (development). It owns the operational semantics of talking to a Kneo Agent Platform instance over its `/v1` HTTP API so that downstream products do not each re-invent them.

Three layers, strict dependency direction *generated* → *core* → *adapters*:

- `kneo_client._generated` — *private*. Generated from the pinned `kneo_serv` OpenAPI spec via `openapi-python-client`. External code must not import from this subpackage.
- `kneo_client.core` — handwritten cross-cutting layer: `Transport`, `Profile`, `ApiKeyAuth`, `RetryPolicy`, idempotency, pagination, request IDs, redacted logging, normalized errors.
- `kneo_client.platform` and `kneo_client.agent` — domain-shaped adapters built on `Transport`. `PlatformClient` for the Dashboard, `AgentClient` for the Studio.

Public entry point is a unified `KneoClient` with `.platform` and `.agent` namespaces backed by a single shared `Transport`.

Privacy boundary. Anything under `kneo_client._generated` is private implementation. Its surface can change between any two `kneo-client` releases without notice. If a generated type or helper is useful to external consumers, it gets re-exported through `core`, `platform`, or `agent`; otherwise it stays internal.

For the architectural rationale see `docs/dev/architecture.md` on the repo. For each major design choice see the ADRs under `docs/dev/adrs/`.

KneoClient

The public entry point. Owns one `Transport` and mounts `.platform` and `.agent` namespaces backed by it. Async context manager; closes its transport on exit.

class KneoClient CLASS

Async client for a Kneo Agent Platform instance.

Attributes

ATTRIBUTE	TYPE	DESCRIPTION
<code>platform</code>	PlatformClient	Operational surface — runs, traces, audit, credentials, policies, health.
<code>agent</code>	AgentClient	Development surface — spec validate / compile / explain / policy_report / run.
<code>profile</code>	Profile	The resolved profile this client is bound to.

```
def __init__(profile: Profile, *, http_client: httpx.AsyncClient | None = None, retry_policy: RetryPolicy | None = None) → None
```

Build a `KneoClient` bound to `profile`. The transport is constructed eagerly; nothing happens on the network until the first `.platform.*` or `.agent.*` call.

PARAMETER	TYPE	DEFAULT	DESCRIPTION
<code>profile</code>	Profile	—	Resolved profile carrying URL, API key, scheme, and timeout.
<code>http_client</code>	httpx.AsyncClient None	None	Pre-built <code>httpx</code> client for custom timeouts / TLS / proxies. Caller owns its lifecycle (not closed on exit); base URL + auth are not overridden.
<code>retry_policy</code>	RetryPolicy None	None	Override the default retry policy (3 attempts, exp backoff). See RetryPolicy .

```
@classmethod def from_profile(name: str | None = None, *, config_file: Path | None = None, url: str | None = None, api_key: str | None = None, auth_scheme: AuthScheme | str | None = None, timeout: float | None = None, http_client: httpx.AsyncClient | None = None) → KneoClient
```

Resolve a profile via `load_profile()` and build a `KneoClient`. The override keywords mirror `load_profile` exactly (fully typed — no `**kwargs`), so callers keep static checking and completion.

PARAMETER	TYPE	DEFAULT	DESCRIPTION
name	str None	None	Profile name to load. Falls back to \$KNEO_PROFILE then "default".
config_file / url / api_key / auth_scheme / timeout	keyword-only	None	Typed overrides forwarded to load_profile (mirrors its signature exactly).
http_client	httpx.AsyncClient None	None	Forwarded to the constructor; caller owns its lifecycle.

```
async def aclose() → None
```

Close the underlying transport. Idempotent.

```
async def __aenter__() → KneoClient · async def __aexit__(*exc) → None
```

Async context manager. Use `async with KneoClient.from_profile() as client:` to ensure the transport closes on exit.

Example

```

import asyncio
from kneo_client import KneoClient

async def main():
    async with KneoClient.from_profile() as client:
        ready = await client.platform.health.readyz()
        print(f"ok={ready.ok}")

        run = await client.platform.runs.create({"spec_id": "my-spec"})
        terminal = await client.platform.runs.wait_for_completion(run.run_id,
            timeout=120)
        print(f"final={terminal.status}")

asyncio.run(main())

```

Transport

The request engine. `Transport` is the only layer that does I/O; every adapter call eventually routes through its `request()` method.

Owens: the `httpx.AsyncClient`, the auth flow, the retry loop, idempotency-key injection, request-ID injection, redacted logging, and the error-to-exception mapping point.

class Transport CLASS

Async HTTP transport for a Kneo Agent Platform instance.

Attributes

ATTRIBUTE	TYPE	DESCRIPTION
<code>profile</code>	Profile	The profile this transport is bound to (read-only).

```

def __init__(profile: Profile, *, http_client: httpx.AsyncClient | None =
None, retry_policy: RetryPolicy | None = None) → None

```

PARAMETER	TYPE	DEFAULT	DESCRIPTION
<code>profile</code>	<code>Profile</code>	—	Resolved profile.
<code>http_client</code>	<code>httpx.AsyncClient</code> <code>None</code>	<code>None</code>	Optional pre-built client. When supplied, the caller owns its lifecycle; <code>aclose()</code> will not close it. Base URL and auth are <i>not</i> overridden on a passed-in client.
<code>retry_policy</code>	<code>RetryPolicy</code> <code>None</code>	<code>None</code>	Defaults to <code>RetryPolicy()</code> (3 attempts, exp backoff with jitter).

```

async def request(method: str, path: str, *, json: Any = None, params: Mapping
| None = None, headers: Mapping | None = None, idempotency_key: str | None =
None, request_id: str | None = None) → httpx.Response

```

Send an HTTP request with auth, retries, idempotency, and error mapping.

PARAMETER	TYPE	DEFAULT	DESCRIPTION
<code>method</code>	<code>str</code>	—	HTTP method (GET , POST , ...).
<code>path</code>	<code>str</code>	—	Path relative to the profile's base URL (e.g. /v1/runs).
<code>json</code>	<code>Any</code>	<code>None</code>	JSON body.
<code>params</code>	<code>Mapping[str, Any]</code> <code>None</code>	<code>None</code>	Query string parameters.
<code>headers</code>	<code>Mapping[str, str]</code> <code>None</code>	<code>None</code>	Extra headers. Auth, Idempotency-Key , and X-Request-ID are added automatically and

			override anything in this mapping.
<code>idempotency_key</code>	<code>str None</code>	<code>None</code>	Override the auto-generated <code>POST</code> idempotency key. Validated against the 256-character platform limit. Ignored on non- <code>POST</code> methods.
<code>request_id</code>	<code>str None</code>	<code>None</code>	Override the auto-generated request ID.

RETURNS

`httpx.Response` — the successful response (status < 400).

RAISES

A [KneoError](#) subclass after retries are exhausted. See [Exceptions](#) for the mapping.

Retry behavior

The transport retries on:

- Transport-level errors from `httpx` (DNS, connect, read timeout, TLS).
- HTTP status in `RETRYABLE_STATUS_CODES = {429, 502, 503, 504}`.

Retries fire only for:

- **Idempotent verbs** — `GET`, `HEAD`, `OPTIONS`, `PUT`, `DELETE`.
- **`POST` with an idempotency key** — which is always, since the transport auto-injects a UUID4 key on every `POST`.

A `Retry-After` response header on 429 / 503 (and the auto-retried 409

`idempotency_key_in_progress`) overrides the computed delay — no jitter on top — when the hint is within `max_delay`. A hint *above* `max_delay` is not slept on: retrying earlier than the server asked is likely doomed, so the transport stops retrying and raises the typed error carrying the verbatim hint on `.retry_after` for caller-side back-off.

```
async def aclose() → None
```

Close the underlying `httpx.AsyncClient` if this `Transport` owns it. When a caller-supplied `http_client` was passed to `__init__`, the caller's client is not closed.

```
async def __aenter__() · async def __aexit__(*exc)
```

Async context manager. `async with Transport(profile) as t:` closes the transport on exit.

Example: direct use (no adapter)

```
import asyncio
from kneo_client.core.profiles import load_profile
from kneo_client.core.transport import Transport

async def main():
    profile = load_profile()
    async with Transport(profile) as t:
        resp = await t.request("GET", "/v1/healthz")
        print(resp.json())

asyncio.run(main())
```

SyncTransport

Synchronous facade around [Transport](#). Runs the async transport in a dedicated background thread + event loop via `anyio.from_thread.start_blocking_portal()`. Each call dispatches into that loop and blocks until the coroutine returns.

For callers that cannot run an event loop (scripts, notebooks, sync frameworks). The async surface remains the recommended path.

class SyncTransport CLASS

```
def __init__(profile: Profile, *, retry_policy: RetryPolicy | None = None) → None
```

```
def request(method: str, path: str, *, json=None, params=None, headers=None, idempotency_key=None, request_id=None) → httpx.Response
```

Synchronous version of `Transport.request`. Same arguments, same semantics; blocks until done.

```
def close() → None · def __enter__() · def __exit__(*exc)
```

Lifecycle. Use with `SyncTransport(profile)` as `t`: to close the background portal on exit.

Caveat. `SyncTransport` does **not** mount `.platform` / `.agent` adapters. Those are async-only. Sync consumers wanting wrapped endpoints either glue async-to-sync at the call site, or drop to `t.request(method, path, ...)` directly.

Example

```
from kneo_client.core.profiles import load_profile
from kneo_client.core.transport import SyncTransport

with SyncTransport(load_profile()) as t:
    resp = t.request("GET", "/v1/healthz")
    print(resp.json())
```

Profile & profiles

A *profile* is a frozen dataclass bundling `(name, url, api_key, auth_scheme, timeout)` — everything `Transport` needs to talk to one Kneo Agent Platform instance.

class Profile DATACLASS

`frozen=True`. Pass directly to `KneoClient(profile)` or build via `load_profile()`.

FIELD	TYPE	DEFAULT	DESCRIPTION
<code>name</code>	<code>str</code>	—	Profile name. Informational; used in log records.
<code>url</code>	<code>str</code>	—	Base URL of the platform instance.

api_key	str	—	API key to authenticate with. Treat like any other secret; the redaction-aware logger masks it in header output.
auth_scheme	AuthScheme	BEARER	How to present the key (Authorization: Bearer vs X-Kneo-Api-Key).
timeout	float	30.0	Per-request timeout in seconds.

class ProfileError EXCEPTION

Raised when a profile cannot be resolved (missing url / api_key , malformed TOML, bad scheme, non-numeric timeout, etc.).

load_profile() FUNCTION

```
def load_profile(name: str | None = None, *, config_file: Path | None = None, url: str | None = None, api_key: str | None = None, auth_scheme: AuthScheme | str | None = None, timeout: float | None = None) → Profile
```

Resolve a Profile from TOML config + env vars + explicit kwargs. Resolution order (later overrides earlier):

- 1. TOML at config_file (default: ~/.config/kneo/client.toml via platformdirs).
- 2. Environment variables: KNEO_URL , KNEO_API_KEY , KNEO_AUTH_SCHEME , KNEO_TIMEOUT .
KNEO_PROFILE selects which TOML section to load.
- 3. Explicit keyword arguments to this function.

PARAMETER	TYPE	DEFAULT	DESCRIPTION
name	str None	None	Profile name. Falls back to \$KNEO_PROFILE , then "default" .

<code>config_file</code>	<code>Path</code> <code>None</code>	<code>None</code>	Override the default TOML location. Non-existent files are skipped silently.
<code>url</code> , <code>api_key</code> , <code>auth_scheme</code> , <code>timeout</code>	<code>str</code> / <code>AuthScheme</code> / <code>float</code>	<code>None</code>	Override the resolved field. <code>auth_scheme</code> accepts either an enum or its string value.

RETURNS

`Profile` — fully resolved.

RAISES

`ProfileError` — if `url` or `api_key` cannot be resolved from any source, or if the config file is malformed.

```
def default_config_path() → Path
```

Return the XDG-style default location of `client.toml` (`~/.config/kneo/client.toml` on Linux, `~/Library/Application Support/kneo/client.toml` on macOS, etc., via `platformdirs.user_config_dir("kneo")`).

Environment-variable constants

CONSTANT	VALUE	PURPOSE
<code>DEFAULT_PROFILE_NAME</code>	<code>"default"</code>	Profile name when none is specified.
<code>DEFAULT_TIMEOUT_SECONDS</code>	<code>30.0</code>	Default per-request timeout.
<code>ENV_PROFILE</code>	<code>"KNEO_PROFILE"</code>	Selects which TOML section to load.
<code>ENV_URL</code>	<code>"KNEO_URL"</code>	Override profile <code>url</code> .
<code>ENV_API_KEY</code>	<code>"KNEO_API_KEY"</code>	Override profile <code>api_key</code> .
<code>ENV_AUTH_SCHEME</code>	<code>"KNEO_AUTH_SCHEME"</code>	Override profile <code>auth_scheme</code> .
<code>ENV_TIMEOUT</code>	<code>"KNEO_TIMEOUT"</code>	Override profile <code>timeout</code> .

Example: TOML config

```
# ~/.config/kneo/client.toml
[default]
url = "https://kneo.example.com"
api_key = "prod-key"
auth_scheme = "bearer"
timeout = 30.0

[staging]
url = "https://staging-kneo.example.com"
api_key = "staging-key"
```

Example: resolve a profile

```
from kneo_client.core.profiles import load_profile

p = load_profile()                # 'default' from TOML + env
p = load_profile("staging")       # explicit profile
p = load_profile(url="https://ad-hoc", api_key=tok) # explicit kwargs win
```

ApiKeyAuth & AuthScheme

The platform accepts the API key as either `Authorization: Bearer <key>` or `X-Kneo-API-Key: <key>`. `ApiKeyAuth` is an `httpx.Auth` subclass that injects the chosen header on every outgoing request; `AuthScheme` is the enum that picks which.

class AuthScheme ENUM

String-valued enum.

MEMBER	VALUE	HEADER INJECTED
BEARER	"bearer"	Authorization: Bearer <key>
KNEO_API_KEY	"kneo_api_key"	X-Kneo-API-Key: <key>

class ApiKeyAuth CLASS

Inherits from `httpx.Auth` . Compatible with both `httpx.Client` and `httpx.AsyncClient` because `httpx`'s auth flow is a synchronous generator.

```
def __init__(api_key: str, scheme: AuthScheme = AuthScheme.BEARER) → None
```

PARAMETER	TYPE	DEFAULT	DESCRIPTION
api_key	str	—	Non-empty API key. Empty values raise <code>ValueError</code> .

scheme	AuthScheme	BEARER	Header scheme.
--------	------------	--------	----------------

```
def auth_flow(request: httpx.Request) → Generator[httpx.Request, httpx.Response, None]
```

Injects the API key header and yields the request unchanged. Called by `httpx` for every redirect / retry attempt at the transport-level layer.

RetryPolicy

A frozen dataclass describing when and how long to wait between attempts. `Transport` applies it; the policy itself does no I/O.

class RetryPolicy DATACLASS

`frozen=True` .

FIELD	TYPE	DEFAULT	DESCRIPTION
max_attempts	int	3	Total attempts including the first try. Set to 1 to disable retries.
base_delay	float	0.2	Seconds to wait before the second attempt.

Each subsequent attempt doubles up to `max_delay` .

<code>max_delay</code>	<code>float</code>	<code>30.0</code>	Cap on the computed delay before jitter.
<code>jitter</code>	<code>float</code>	<code>0.1</code>	Fraction of the delay to randomize by, in <code>[0, 1]</code> . With <code>jitter=0.1</code> and a 1-second delay, the actual sleep is uniformly in <code>[0.9, 1.1]</code> .

```
def delay_for(attempt: int, retry_after: float | None = None) → float
```

Compute the sleep duration before `attempt` (1-indexed). `attempt=1` returns 0 (no delay before the first try).

PARAMETER	TYPE	DEFAULT	DESCRIPTION
<code>attempt</code>	<code>int</code>	<code>—</code>	The upcoming attempt number, starting at 2.
<code>retry_after</code>	<code>float None</code>	<code>None</code>	Optional server-supplied hint (from a <code>Retry-After</code> header). Honored verbatim — no jitter — up to <code>max_delay</code> . Hints above the cap never reach this method: the transport fails fast on them, surfacing the typed error with <code>.retry_after</code> instead of retrying earlier than the server asked; the <code>min</code> clamp

here is defense in
depth for direct callers.

RETRYABLE_STATUS_CODES

Module constant: frozenset({429, 502, 503, 504}) . Intentionally not configurable; if your platform deployment legitimately returns transient 500 s, fix the platform.

Example

```
from kneo_client import KneoClient
from kneo_client.core.retries import RetryPolicy

# Aggressive policy for a flaky network
client = KneoClient(profile, retry_policy=RetryPolicy(max_attempts=8,
base_delay=0.5))

# No retries (first failure surfaces immediately)
client = KneoClient(profile, retry_policy=RetryPolicy(max_attempts=1))
```

Idempotency helpers

Every POST the client sends carries an Idempotency-Key . Transport auto-injects a fresh UUID4 per request unless the caller supplies one via the method's idempotency_key= parameter.

The platform short-circuits a duplicate POST with the same key + identical payload, returning the original response. Reuse of the same key with a different payload yields HTTP 409, surfaced as [KneoIdempotencyMismatchError](#) .

CONSTANT	VALUE	PURPOSE
IDEMPOTENCY_KEY_HEADER	"Idempotency-Key"	The header name as the platform expects it.
MAX_KEY_LENGTH	256	Platform-enforced maximum length for the value.

```
def new_idempotency_key() → str
```

Return a fresh UUID4-based idempotency key as a string.

```
def validate_idempotency_key(key: str) → None
```

Validate an externally-supplied key. Raises `ValueError` if empty or longer than `MAX_KEY_LENGTH` .

List-method results

Module `kneo_client.core.results` (renamed from `core.pagination` in 0.2.0). Hosts the two collection wrappers every list-method on `PlatformClient` returns, plus the page-walking `iterate_all()` .

class Page[T] DATACLASS

Wrapper around list-shaped server responses. Every paginating endpoint populates the full metadata set against current servers (audit included — fully paginated since `kneo_serv 0.6.0` / client 0.6.0). Metadata fields are all Optional: against an older server that omits a field, the field surfaces as `None` rather than carrying a synthesized value. Supports sequence ergonomics: `__iter__` , `__len__` , `__getitem__` .

FIELD	TYPE	DESCRIPTION
<code>items</code>	<code>list[T]</code>	Items on this page.
<code>total</code>	<code>int</code> <code>None</code>	Total items across all pages, when the server echoes it.
<code>limit</code>	<code>int</code> <code>None</code>	Page size the server applied, when echoed.
<code>offset</code>	<code>int</code> <code>None</code>	Offset of the first item on this page, when echoed.
<code>sort_by</code>	<code>str</code> <code>None</code>	Sort field in effect, when echoed.
<code>sort_order</code>	<code>str</code> <code>None</code>	"asc" or "desc" , when echoed.

<code>window</code>	<code>int None</code>	The server's paging window — the deepest offset it will page to (<code>kneo_serv >= 0.9.0</code>). <code>total</code> is the true store count and can exceed it; <code>has_more</code> accounts for both.
---------------------	-------------------------	---

Properties

PROPERTY	TYPE	DESCRIPTION
<code>count</code>	<code>int</code>	<code>len(items)</code> .
<code>has_more</code>	<code>bool</code>	<code>offset + count < min(total, window)</code> — reports whether more items are <i>fetchable</i> , clamping to the paging window when the server discloses one, so pagination loops stop honestly at the window edge instead of fetching one guaranteed-empty page. <code>False</code> when <code>offset</code> or <code>total</code> is <code>None</code> (server didn't tell us, so we don't speculate).

class Map[K, V] DATACLASS

Wrapper around dict-keyed server responses (`credentials.list` , `policies.environment_list`). The whole collection comes back in one call; there is no pagination. Exposes dict-style accessors.

FIELD	TYPE	DESCRIPTION
<code>items</code>	<code>Mapping[K, V]</code>	The underlying keyed collection.

Properties + methods

MEMBER	TYPE	DESCRIPTION
<code>count</code>	<code>int</code>	<code>len(items)</code> .
<code>__getitem__</code>	<code>(K) → V</code>	Bracket lookup.

<code>__iter__</code>	<code>→ Iterator[V]</code>	Iterates <i>values</i> (consistent with <code>Page</code> iterating items).
<code>__contains__</code>	<code>(object) → bool</code>	Key-based membership.
<code>get</code>	<code>(K, default=None) → V None</code>	Safe lookup.
<code>keys</code>	<code>() → KeysView[K]</code>	Keys view.
<code>values</code>	<code>() → ValuesView[V]</code>	Values view.

```
async def iterate_all(fetch_page: Callable[[int, int], Awaitable[Page[T]]], *,
page_size: int = 100, start_offset: int = 0) → AsyncIterator[T]
```

Walk all items across pages. `fetch_page(limit, offset)` must return a `Page`. Clamps `page_size` to `MAX_PAGE_SIZE = 1000`. Stops when `Page.has_more` is `False` — the server signaled end-of-data, or the page reached the paging window edge. Works across every paginating endpoint, audit included (`audit.list` is fully paginated since client 0.6.0). Against a pre-0.6.0 *server* that omits the pagination metadata, the fields degrade to `None` and the walk does one fetch then exits.

Constants: `DEFAULT_PAGE_SIZE = 100` , `MAX_PAGE_SIZE = 1000` .

Request-ID helpers

Every request carries an `X-Request-ID` . The transport auto-injects a UUID4; callers can override via the method's `request_id=` parameter for cross-service correlation. The platform echoes the ID on responses and in audit events.

CONSTANT	VALUE
<code>REQUEST_ID_HEADER</code>	<code>"X-Request-ID"</code>

```
def generate_request_id() → str
```

Return a fresh UUID4-based request ID as a string.

Logging helpers

Loggers under the `kneo_client.*` namespace use the standard logging machinery. All header payloads are passed through `redact_headers()` before any sink sees them.

CONSTANT	VALUE
<code>REDACTED</code>	<code>"<redacted>"</code>

```
def get_logger(name: str = "kneo_client") → logging.Logger
```

Return a logger under the `kneo_client.*` hierarchy. Bare names like `"transport"` are namespaced to `kneo_client.transport`; already-namespaced names (e.g. `"kneo_client.foo"`) pass through unchanged.

```
def redact_headers(headers: Mapping[str, str]) → dict[str, str]
```

Return a copy of `headers` with sensitive values replaced by `<redacted>`. Sensitive header tokens (case-insensitive, substring match): `authorization`, `x-kneo-api-key`, `cookie`, `set-cookie`, `proxy-authorization`.

Body redaction is the caller's responsibility. The redaction layer masks headers, not bodies. If you log request / response payloads, filter them at your application's logging-formatter layer.

Exceptions

Every failure surfaces as a typed exception derived from `KneoError`. Each one carries the original HTTP status, response body (parsed JSON when possible), the server-assigned request ID, and — for `POST` failures — the idempotency key that was sent.

Hierarchy

```
KneoError |— KneoNetworkError # transport-level: DNS / connect / TLS / read
timeout |— KneoBadRequestError # HTTP 400 – semantically rejected (branch on
.code) |— KneoAuthError # HTTP 401 |— KneoPermissionError # HTTP 403 |—
KneoNotFoundError # HTTP 404 |— KneoConflictError # HTTP 409 (carries
```

```
.retry_after; branch on .code) | └─ KneoIdempotencyMismatchError # HTTP 409
with code idempotency_key_conflict |─ KneoPayloadTooLarge # HTTP 413 (carries
.max_body_bytes) |─ KneoValidationError # HTTP 422 – request validation failed
|─ KneoRateLimited # HTTP 429 (carries .retry_after) |─ KneoServerError # HTTP
5xx |─ KneoServiceUnavailable # HTTP 503 – backpressure / overload (carries
.retry_after)
```

class KneoError EXCEPTION

Base exception. Inherits from `Exception`.

ATTRIBUTE	TYPE	DESCRIPTION
<code>status</code>	<code>int None</code>	HTTP status, or <code>None</code> for transport-level failures.
<code>body</code>	<code>Any</code>	Parsed JSON dict, raw text, or <code>None</code> .
<code>request_id</code>	<code>str None</code>	The X-Request-ID the platform returned, when present.
<code>idempotency_key</code>	<code>str None</code>	The Idempotency-Key sent on the failing request, when set.
<code>code</code>	<code>str None</code>	Stable snake_case error code from the platform's error envelope (<code>detail.error; kneo_serv 0.6.0+</code>), or <code>None</code> when the body carries no code. Branch on this rather than the human-readable message.

class KneoBadRequestError EXCEPTION

HTTP 400 — the request was structurally valid but semantically rejected. The platform's stable codes include `invalid_request`, `spec_invalid` (with a diagnostics list in the body), and `token_budget_exceeded` — branch on `KneoError.code`.

class KneoConflictError EXCEPTION

HTTP 409 — a conflict with current server state. The platform's 409s carry a stable `code` discriminating the cause: `run_state_conflict` (a lifecycle fence, e.g. cancelling an already-terminal run or resuming a run that isn't blocked) and `resource_locked` (held by another operation) raise this class; `idempotency_key_in_progress` (a request with the same `Idempotency-Key` is still executing) is transient — the transport auto-retries it honoring `Retry-After` , so it reaches the caller as this class with `retry_after` set only once retries are exhausted; `idempotency_key_conflict` raises the `KneoIdempotencyMismatchError` subclass instead. Adds one attribute on top of `KneoError` :

ATTRIBUTE	TYPE	DESCRIPTION
<code>retry_after</code>	<code>float</code> <code>None</code>	Seconds parsed from the <code>Retry-After</code> header, or <code>None</code> when the server did not provide the hint.

class KneoIdempotencyMismatchError EXCEPTION

HTTP 409 caused by an `Idempotency-Key` replay with a different payload — almost always a caller bug. A subclass of `KneoConflictError` . Raised only for the `idempotency_key_conflict` envelope code (or, against pre-0.6.0 servers whose bodies carry no code, for any 409 on a keyed request); other 409 causes raise the `KneoConflictError` base instead.

class KneoPayloadTooLarge EXCEPTION

HTTP 413 — the request body exceeds the server's size cap (`kneo_serv` enforces `KNEO_SERV_MAX_BODY_BYTES` , chunked bodies included). Not retryable — shrink the payload. Adds one attribute on top of `KneoError` :

ATTRIBUTE	TYPE	DESCRIPTION
<code>max_body_bytes</code>	<code>int</code> <code>None</code>	The server's configured cap, when disclosed in the error envelope.

class KneoValidationError EXCEPTION

HTTP 422 — request validation failed. Two body shapes reach this: the platform's error envelope (stable codes such as `guardrail_violation`), and FastAPI's list-shaped request-validation detail (`{"detail":`

[{loc, msg, type}, ...] }, whose first entry is summarized into the exception message. The full diagnostics stay available on `KneoError.body`.

class KneoRateLimited EXCEPTION

Adds one attribute on top of `KneoError` :

ATTRIBUTE	TYPE	DESCRIPTION
retry_after	float None	Seconds parsed from the Retry-After header.

class KneoServiceUnavailable EXCEPTION

HTTP 503 — the platform is temporarily unavailable or shedding load (notably `kneo_serv` 's run-queue backpressure on `POST /v1/runs`). A **subclass of `KneoServerError`** , so except `KneoServerError` still catches it. The transport already auto-retries 503 (honoring `Retry-After`); this surfaces only once retries are exhausted. Adds one attribute on top of `KneoError` :

ATTRIBUTE	TYPE	DESCRIPTION
retry_after	float None	Seconds parsed from the Retry-After header, or None when the server did not provide the hint.

from_response() FUNCTION

```
def from_response(response: httpx.Response, *, idempotency_key: str | None = None) → KneoError
```

Map an HTTP response to the appropriate `KneoError` subclass.

HTTP status	Exception	Notes
400	KneoBadRequestError	Semantically rejected; branch on <code>code</code> (<code>invalid_request</code> , <code>spec_invalid</code> , <code>token_budget_exceeded</code>).
401	KneoAuthError	Missing or invalid API key.

403	KneoPermissionError	Key valid; insufficient scope.
404	KneoNotFoundError	Resource missing.
409, code idempotency_key_conflict	KneoIdempotencyMismatchError	Idempotency-key replay with different payload. Also raised for code-less legacy bodies (pre-0.6.0 servers) on keyed requests.
409 (other codes)	KneoConflictError	run_state_conflict / resource_locked ; idempotency_key_in_progress is auto-retried by the transport and surfaces here with retry_after once exhausted.
413	KneoPayloadTooLarge	Body exceeds the server's size cap; carries max_body_bytes .
422	KneoValidationError	Envelope codes (e.g. guardrail_violation) or FastAPI list-shaped detail, summarized into the message.
429	KneoRateLimited	Carries retry_after .
503	KneoServiceUnavailable	Backpressure / overload. KneoServerError subclass; carries retry_after .
5xx (other)	KneoServerError	Server-side failure.
Other	KneoError	Catch-all.
Transport	KneoNetworkError	Wrapped from httpx.HTTPError .

Example

```
from kneo_client.core.errors import KneoError, KneoIdempotencyMismatchError,
KneoRateLimited

try:
    run = await client.platform.runs.create(payload, idempotency_key=key)
except KneoIdempotencyMismatchError as exc:
    # Same key reused with a different payload – almost always a caller bug.
    log.error("mismatch on key=%r body=%r", exc.idempotency_key, exc.body)
    raise
except KneoRateLimited as exc:
    await asyncio.sleep(exc.retry_after or 10)
except KneoError as exc:
    log.error("create_run failed status=%s rid=%s", exc.status, exc.request_id)
    raise
```

PlatformClient

Operational surface for **Kneo Agent Dashboard**. Aggregates six sub-clients backed by a shared [Transport](#).

class PlatformClient CLASS

ATTRIBUTE	TYPE	DESCRIPTION
health	HealthClient	/v1/{healthz,livez,readyz }
runs	RunsClient	/v1/runs + 11 sub-endpoints
human_tasks	HumanTasksClient	/v1/human-tasks family
audit	AuditClient	/v1/audit-events
credentials	CredentialsClient	/v1/security/credentials
policies	PoliciesClient	/v1/policies/environment family

```
def __init__(transport: Transport) → None
```

Build the aggregator and attach all six sub-clients. Usually you don't construct one directly — `KneoClient(profile).platform` gives you one.

HealthClient

Wraps the three platform health probes.

class HealthClient CLASS

```
async def healthz() → HealthResponse
```

GET `/v1/healthz` — overall service health.

```
async def livez() → HealthResponse
```

GET `/v1/livez` — process liveness. Does *not* check downstream deps.

```
async def readyz() → HealthResponse
```

GET `/v1/readyz` — readiness (database, queue, runtime registry, providers, MCP).

Response (HealthResponse)

FIELD	TYPE	DESCRIPTION
<code>ok</code>	<code>bool</code>	Overall pass/fail.
<code>service</code>	<code>str</code> <code>UNSET</code>	Service name, e.g. "kneo-serv-platform" .
<code>version</code>	<code>str</code> <code>UNSET</code>	Platform version.
<code>metadata</code>	<code>HealthResponseMetadata</code> <code>UNSET</code>	Per-check breakdown.

RunsClient

The largest sub-client — 12 endpoints + one convenience helper (`wait_for_completion`). Covers the full run lifecycle: create, list, get, cancel, continue, replay, trace, checkpoints, policy reports, recovery.

class RunsClient CLASS

```
async def create(body: RunCreateRequest | dict, *, idempotency_key: str | None = None) → RunCreateResponse
```

POST /v1/runs — start a new run. Idempotency key is auto-generated unless caller supplies one.

```
async def list(*, status: str | None = None, limit: int = 100, offset: int = 0, sort_by: str = "updated_at", sort_order: str = "desc") → Page[RunListResponseRunsItem]
```

GET /v1/runs — list runs with limit/offset pagination.

```
async def get(run_id: str) → RunStatusResponse
```

GET /v1/runs/{run_id} — fetch the current status of a run.

```
async def cancel(run_id: str, *, idempotency_key: str | None = None) → RunStatusResponse
```

POST /v1/runs/{run_id}/cancel — cooperatively cancel a running run. The platform may not honor the cancel if the run is already in a terminal state; the returned status reflects the post-cancel state.

```
async def continue_(run_id: str, *, idempotency_key: str | None = None) → RunCreateResponse
```

POST /v1/runs/{run_id}/continue — resume a paused run (typically after a human task was resumed). Trailing underscore avoids the Python `continue` keyword collision. The auto-injected `Idempotency-Key` makes replays contractually safe on `kneo_serv >= 0.9.0` (older servers treated the key as best-effort on this route).

```
async def replay(run_id: str) → RunReplayResponse
```

GET /v1/runs/{run_id}/replay — fetch a deterministic replay view derived from the run's checkpoints + trace.

```
async def trace(run_id: str, *, event_type: str | None = None, limit: int = 100, offset: int = 0, sort_by: str = "timestamp", sort_order: str = "asc") → Page[TraceResponseEventsItem]
```

GET /v1/runs/{run_id}/trace — fetch the event trace for a run (tool calls, model calls, middleware decisions, etc.).

```
async def checkpoints(run_id: str, *, type: str | None = None, limit: int = 100, offset: int = 0, sort_by: str = "sequence", sort_order: str = "asc") → Page[CheckpointListResponseCheckpointsItem]
```

GET /v1/runs/{run_id}/checkpoints — list a run's serialized state snapshots.

```
async def checkpoints_diff(run_id: str, *, from_sequence: int | None = None, to_sequence: int | None = None) → CheckpointDiffResponse
```

GET /v1/runs/{run_id}/checkpoints/diff — diff two checkpoints.

```
async def policy_report(run_id: str) → SpecPolicyReportResponse
```

GET /v1/runs/{run_id}/policy-report — fetch policy outcomes for a run. Returns the typed `SpecPolicyReportResponse`; the spec types this 200 response identically to `POST /v1/specs/policy-report`, so it mirrors [SpecsClient.policy_report](#).

```
async def recovery(run_id: str) → RunRecoveryResponse
```

GET /v1/runs/{run_id}/recovery — fetch the recovery context for a failed run (what state can be recovered, what next action is recommended).

```
async def wait_for_completion(run_id: str, *, timeout: float | None = None, poll_interval: float = 1.0, terminal_statuses: Iterable[str] | None = None) → RunStatusResponse
```

Poll `get(run_id)` until the run reaches a terminal status. Convenience helper; not a separate platform endpoint.

PARAMETER	TYPE	DEFAULT	DESCRIPTION
-----------	------	---------	-------------

<code>timeout</code>	<code>float None</code>	<code>None</code>	Total time budget in seconds. <code>None</code> means wait indefinitely.
<code>poll_interval</code>	<code>float</code>	<code>1.0</code>	Seconds to sleep between polls.
<code>terminal_statuses</code>	<code>Iterable[str] None</code>	<code>None</code>	Defaults to <code>{"completed", "failed", "cancelled", "timed_out", "expired"}</code> . Pass an explicit set including <code>"blocked"</code> to stop waiting when the run pauses for human review.

RETURNS

The terminal `RunStatusResponse` .

RAISES

`TimeoutError` if `timeout` elapses before a terminal status is reached.

Module constant: `DEFAULT_TERMINAL_STATUSES = frozenset({"completed", "failed", "cancelled", "timed_out", "expired"})` .

```
async def tail_trace(run_id: str, *, poll_interval: float = 1.0, timeout: float | None = None, terminal_statuses: Iterable[str] | None = None, start_offset: int = 0, page_size: int = 100) → AsyncIterator[TraceResponseEventsItem]
```

Stream trace events as they arrive. Polls `trace(run_id)` with ascending offset, yields each new event as the page arrives, then performs one final drain pass after the run reaches a terminal status to capture any events emitted between the last poll and the status transition. Returns when the drain pass finds no new events. Convenience helper; not a separate platform endpoint.

PARAMETER	TYPE	DEFAULT	DESCRIPTION
-----------	------	---------	-------------

<code>poll_interval</code>	<code>float</code>	<code>1.0</code>	Seconds to sleep between polls.
<code>timeout</code>	<code>float None</code>	<code>None</code>	Total time budget. <code>None</code> means wait indefinitely.
<code>terminal_statuses</code>	<code>Iterable[str] None</code>	<code>None</code>	Status set that ends streaming. Defaults to <code>{"completed", "failed", "cancelled", "timed_out", "expired"}</code> ; include <code>"blocked"</code> to stop at a human-review pause.
<code>start_offset</code>	<code>int</code>	<code>0</code>	Offset to start from. Use a saved checkpoint to resume mid-trace.
<code>page_size</code>	<code>int</code>	<code>100</code>	Per-poll page size; clamped server-side to 1000.

RAISES

`TimeoutError` if `timeout` elapses before a terminal status is reached.

When to pick this over `wait_for_completion` — see [polling and waiting guide](#). Short version:

`wait_for_completion` when you only need the final status; `tail_trace` when you need to display, log, or react to events as they arrive (typically interactive UIs and dev tooling).

Example: end-to-end run with live trace

```

async with KneoClient.from_profile() as client:
    created = await client.platform.runs.create({"spec_id": "my-spec"})
    async for ev in client.platform.runs.tail_trace(created.run_id, timeout=600):
        print(f"{ev['sequence']}: {ev['event_type']}")
    # tail_trace returned → the run reached a terminal status and the drain pass
    # found no more events.
    terminal = await client.platform.runs.get(created.run_id)
    print(f"final={terminal.status}")

```

HumanTasksClient

Human-in-the-loop pause points. A run that requires operator review surfaces as a *human task*, identified by a `continuation_id`. Resuming posts a decision and unblocks the paused continuation.

class HumanTasksClient CLASS

```

async def list(*, status: str | None = None, run_id: str | None = None,
               workflow_kind: str | None = None, limit: int = 100, offset: int = 0, sort_by:
               str | None = None, sort_order: str | None = None) →
    Page[HumanTaskListResponseTasksItem]

```

GET `/v1/human-tasks` — list pending and recent human tasks as a `Page`. `run_id` and `workflow_kind` are server-side filters (added in 0.2.0). `status` accepts `pending` (awaiting review) or `escalated` (the task's escalation timeout fired); honored by `kneo_serv >= 0.8.0`, silently ignored by older servers. `sort_by` / `sort_order` default to `None` — when unset, the server's own defaults apply.

```

async def get(continuation_id: str) → HumanTaskResponse

```

GET `/v1/human-tasks/{continuation_id}` — fetch a single human task.

```

async def resume(continuation_id: str, body: HumanResumeRequest | dict, *,
                 idempotency_key: str | None = None) → HumanResumeResponse

```

POST `/v1/human-tasks/{continuation_id}/resume` — resume a paused run with a decision.

```
async def resume_by_id(continuation_id: str, body: HumanResumeRequest | dict, *, idempotency_key: str | None = None) → HumanResumeResponse
```

Fetch a task by ID, then resume it. Calls `get(continuation_id)` first so a stale `continuation_id` surfaces as `KneoNotFoundError` from the GET rather than a conflict from the POST. Use this when distinguishing "the task is gone" from "the resume payload conflicts" matters; call `resume()` directly otherwise.

```
async def resume_first_pending(body: HumanResumeRequest | dict, *, run_id: str | None = None, workflow_kind: str | None = None, idempotency_key: str | None = None) → HumanResumeResponse | None
```

List pending tasks matching the filters and resume the first match; returns `None` if no matching pending task exists. Worker-pattern convenience: collapses the list-then-resume dance into one call. The list call uses `limit=1` and `status="pending"`.

PARAMETER	TYPE	DEFAULT	DESCRIPTION
<code>run_id</code>	<code>str None</code>	<code>None</code>	Filter — only consider tasks for this run.
<code>workflow_kind</code>	<code>str None</code>	<code>None</code>	Filter — only consider tasks of this workflow kind.
<code>idempotency_key</code>	<code>str None</code>	<code>None</code>	Optional caller-supplied key for the POST ; transport generates a fresh UUID4 by default.

AuditClient

Audit events are the canonical operational log: every run, human-task, credential, and policy mutation produces one.

class AuditClient CLASS

```

async def list(*, event_type: str | None = None, run_id: str | None = None,
limit: int = 100, offset: int | None = None, sort_by: str | None = None,
sort_order: str | None = None) → Page[AuditEventListResponseEventsItem]

```

GET /v1/audit-events — list audit events as a fully paginated Page . kneo_serv 0.6.0 added server-honored offset / sort_by / sort_order , so the response carries total / offset metadata and iterate_all walks it. *(These kwargs were removed in 0.5.0 as dead — the platform did not honor them then — and re-introduced in 0.6.0 once kneo_serv 0.6.0 made them real. principal stays removed: still not a documented filter. Against an older server that omits the metadata the Page fields degrade to None .)*

CredentialsClient

Lists the credential *references* the platform knows about. The platform never returns raw secret material via the HTTP API.

class CredentialsClient CLASS

```

async def list(*, providers: list[str] | None = None, extras: list[str] | None
= None, include_service_tokens: bool = True) → Map[str, Any]

```

GET /v1/security/credentials — list known credential references as a Map keyed by credential ID. Credential body is open-shaped in the pinned spec, so values are Any . The whole inventory is returned in one call (no pagination). Filter via providers=["aws", "gcp"] , extras=["scope"] , or include_service_tokens=False ; the three kwargs were added in 0.3.0.

PoliciesClient

Environment policies map deployment targets (e.g. dev , staging , prod) to policy bundles that gate which specs / agents are allowed to run there.

class PoliciesClient CLASS

```
async def environment_list() → Map[str,
EnvironmentPolicyListResponsePoliciesAdditionalProperty]
```

GET /v1/policies/environment — list policies for every environment as a Map keyed by environment name. The whole inventory is returned in one call (no pagination).

```
async def environment_get(environment: str) → EnvironmentPolicyResponse
```

GET /v1/policies/environment/{environment} — fetch one environment's policy.

```
async def environment_put(environment: str, body: EnvironmentPolicyRequest |
dict) → EnvironmentPolicyResponse
```

PUT /v1/policies/environment/{environment} — replace an environment's policy.

Note. PUT is idempotent by HTTP semantics; the transport does not inject an Idempotency-Key for it.

AgentClient

Development surface for **Kneo Agent Studio**. Aggregates two sub-clients ([SkillsClient](#) and [SpecsClient](#)) backed by the shared [Transport](#) .

class AgentClient CLASS

ATTRIBUTE	TYPE	DESCRIPTION
skills	SkillsClient	/v1/skills
specs	SpecsClient	/v1/specs/*

SkillsClient

Read-only skill catalog: the declared + default skills a spec (or a per-request SkillsOverlay on [runs.create](#)) may reference by name. The endpoint requires kneo_serv >= 0.8.0 ; older servers

return 404, surfaced as `KneoNotFoundError`. Exposed as `client.agent.skills`.

class SkillsClient CLASS

```
async def list(*, limit: int = 100, offset: int = 0, sort_by: str = "name",
               sort_order: str = "asc") → Page[SkillsListResponseSkillsItem]
```

GET `/v1/skills` — list the skill catalog as a fully paginated `Page` (metadata includes the `window` field on `kneo_serv >= 0.9.0`). Each item is an open mapping carrying at least `name` / `description` / `path`. The names returned here are the valid targets for a `SkillsOverlay.add` list on `runs.create` — the server rejects overlay names not declared in the spec's skills block.

SpecsClient

Spec validation, compilation, explanation, policy-report, and ad-hoc dry-run. The Studio iterate-and-test loop.

class SpecsClient CLASS

```
async def validate(body: SpecValidateRequest | dict, *, idempotency_key: str |
                  None = None) → SpecValidateResponse
```

POST `/v1/specs/validate` — schema + semantic validation of a spec.

```
async def compile(body: SpecCompileRequest | dict, *, idempotency_key: str |
                  None = None) → SpecCompileResponse
```

POST `/v1/specs/compile` — compile a spec to its runtime representation.

```
async def explain(body: SpecExplainRequest | dict, *, idempotency_key: str |
                  None = None) → SpecExplainResponse
```

POST `/v1/specs/explain` — human-readable summary of a spec.

```
async def policy_report(body: SpecPolicyReportRequest | dict, *,
                       idempotency_key: str | None = None) → SpecPolicyReportResponse
```

POST /v1/specs/policy-report — preview policy outcomes for a spec.

```
async def run(body: RunCreateRequest | dict, *, idempotency_key: str | None = None) → RunCreateResponse
```

POST /v1/specs/run — ad-hoc dry-run of a spec. Distinct from `RunsClient.create`: `specs.run` accepts an inline spec rather than a spec reference, and is intended for Studio's iterate-and-test flow.

```
async def validate_then_compile(body: SpecValidateRequest | dict, *, idempotency_key: str | None = None) → ValidateThenCompileResult
```

Validate first, then compile only if validate succeeded. Skips the compile call on validate failure to avoid a wasted round-trip on a known-bad spec; `result.compile` is `None` in that case and `result.ok` returns `False`. Same body shape as `validate / compile`; the helper passes the same payload to both calls.

```
async def dry_run(body: SpecValidateRequest | dict, *, idempotency_key: str | None = None) → DryRunResult
```

Validate, then run `explain` + `policy_report` in parallel via `asyncio.gather` once validate succeeds. Skips both downstream calls on validate failure. Returns a frozen `DryRunResult` with the validate result, the (optional) `explain` and `policy-report` responses, and an `.ok` property.

Example: validate-then-compile via the helper

```
async with KneoClient.from_profile() as client:
    payload = {"spec": spec_text}
    result = await client.agent.specs.validate_then_compile(payload)
    if not result.ok:
        for d in (result.validate.diagnostics or []):
            print(f"validate: {d}")
        if result.compile is not None:
            for d in (result.compile.diagnostics or []):
                print(f"compile: {d}")
    return
# result.compile is the compiled spec; ship it.
```

ValidateThenCompileResult

Frozen dataclass returned by `SpecsClient.validate_then_compile`. Bundles the validate and (optional) compile responses with a single boolean for the common "is the spec shippable?" question.

class ValidateThenCompileResult DATACLASS

FIELD	TYPE	DESCRIPTION
<code>validate</code>	<code>SpecValidateResponse</code>	Always populated.
<code>compile</code>	<code>SpecCompileResponse</code> <code>None</code>	<code>None</code> when validate failed (the helper skipped the compile call).

Properties

PROPERTY	TYPE	DESCRIPTION
<code>ok</code>	<code>bool</code>	True iff <code>validate.valid</code> AND <code>compile</code> is not <code>None</code> AND <code>compile.ok</code> .

DryRunResult

Frozen dataclass returned by `SpecsClient.dry_run`. Bundles validate + explain + policy-report into a single object for Studio's pre-ship dashboard.

class DryRunResult DATACLASS

FIELD	TYPE	DESCRIPTION
<code>validate</code>	<code>SpecValidateResponse</code>	Always populated.
<code>explain</code>	<code>SpecExplainResponse</code> <code>None</code>	<code>None</code> when validate failed.
<code>policy_report</code>	<code>SpecPolicyReportResponse</code> <code>None</code>	<code>None</code> when validate failed.

Properties

PROPERTY	TYPE	DESCRIPTION
<code>ok</code>	<code>bool</code>	True iff validate succeeded AND <code>policy_report</code> is not None AND <code>policy_report.valid</code> is not False. The UNSET case for <code>policy_report.valid</code> is treated as "no objection" (returns True) — absence is not failure.

Compatibility matrix

Which `kneo-client` release supports which `kneo_serv` platform version. The platform's `/v1` HTTP API is a stability boundary.

<code>kneo-client</code>	Pinned to <code>kneo_serv</code>	Tested against	Python	Status
0.7.0	<code>v0.9.0</code> (<code>info.version</code> 0.9.0)	<code>kneo_serv</code> 0.9.x line	≥ 3.12 (3.12 / 3.13 / 3.14 in CI)	Current
0.6.0	<code>v0.7.0</code> (<code>info.version</code> 0.7.0)	<code>kneo_serv</code> 0.7.x line	≥ 3.12 (3.12 / 3.13 / 3.14 in CI)	Previous
0.5.0	<code>v0.5.0</code> (<code>info.version</code> 0.5.0)	<code>kneo_serv</code> 0.5.x line	≥ 3.12 (3.12 / 3.13 / 3.14 in CI)	Older
0.4.0	<code>v0.4.0</code> (<code>info.version</code> 0.4.0)	<code>kneo_serv</code> 0.4.x line	≥ 3.12 (3.12 / 3.13 / 3.14 in CI)	Older

0.3.0	v0.4.0 (info.version 0.4.0)	kneo_serv 0.4.x line	≥ 3.12 (3.12 / 3.13 / 3.14 in CI)	Older
--------------	-------------------------------------	-------------------------	--------------------------------------	-------

Forward compatibility

A newer `kneo_serv` that adds endpoints will still work for everything `kneo-client` already wraps. New endpoints are available via the drop-to-transport escape hatch until the next `kneo-client` release wraps them:

```
async with KneoClient.from_profile() as client:
    resp = await client._transport.request("GET", "/v1/some/new/endpoint")
    payload = resp.json()
```

Backward compatibility

`kneo-client` X.Y.Z is **not** guaranteed against `kneo_serv` releases older than its pin. Wrappers may rely on response fields that older versions don't emit, and error mapping assumes the current platform error shape. Pin to a matching `kneo-client` minor when talking to an older platform.

Guides

Walking pages with `iterate_all()`

List methods return `Page` (or `Map` for the two dict-keyed endpoints) directly — no adapter needed. Pass the list method itself as the `fetch_page` callable:

```
from kneo_client.core.results import iterate_all

async def all_runs(client, **filters):
    async def fetch_page(limit, offset):
        return await client.platform.runs.list(limit=limit, offset=offset,
        **filters)
    async for run in iterate_all(fetch_page, page_size=200):
        yield run
```

This works for every paginating endpoint, audit included — `audit.list` has been fully paginated since client 0.6.0 (`offset` / `sort_by` / `sort_order` kwargs; `total` / `offset` echoed in the `Page`), so `iterate_all` walks all audit pages. (Against a pre-0.6.0 *server* the metadata degrades to `None` and the walk does one fetch then exits.) For the two `Map` -returning endpoints, walk the values directly: `for cred in (await client.platform.credentials.list()): .`

Bring your own httpx client

For shared connection pools, custom transports, proxy configuration:

```
import httpx
from kneo_client.core.auth import ApiKeyAuth, AuthScheme
from kneo_client.core.profiles import Profile
from kneo_client.core.transport import Transport

profile = Profile(name="prod", url="https://kneo", api_key="...",
                  auth_scheme=AuthScheme.BEARER, timeout=30.0)
shared = httpx.AsyncClient(
    base_url=profile.url,
    auth=ApiKeyAuth(profile.api_key, profile.auth_scheme),
    timeout=profile.timeout,
    limits=httpx.Limits(max_keepalive_connections=20, max_connections=100),
)
try:
    async with Transport(profile, http_client=shared) as t:
        # t.aclose() will NOT close `shared`.
        ...
finally:
    await shared.aclose()
```

Wait for a non-default terminal status

`wait_for_completion()` treats `{"completed", "failed", "cancelled", "timed_out", "expired"}` as terminal by default. To return as soon as the run pauses for human review:

```
terminal = await client.platform.runs.wait_for_completion(  
    run_id,  
    poll_interval=2.0,  
    timeout=300,  
    terminal_statuses={"completed", "failed", "cancelled", "timed_out", "expired",  
"blocked"},  
)
```

More recipes

See `docs/dev/extending.md` on the repo for the full 10-recipe set: custom retry policy, custom auth scheme, wrapping a not-yet-adapter-covered endpoint, adding a new wrapped endpoint, custom logging, profile from a secrets manager, paginated iteration, custom terminal statuses, sync facade, and what NOT to extend.

Kneo Agent Client — API Reference Manual · v0.7.0 · 2026-06-12 · MIT license

Adapted layout and style from the `kneo_agent` SDK API reference.