

Kneo Agent Client

User Guide

Version 0.9.0

2026-06-21

Quickstart, profiles and auth, idempotency and retries, pagination, polling and waiting, error handling, and the compatibility matrix — one printable reading path for the kneo-client user documentation.

Table of contents

Table of contents	2
Quickstart	5
What "quickstart" means in kneo-client	5
Install	5
Configure a profile	5
Ping the platform	6
Create your first run	7
Skills catalog and per-request overlays	8
Sync facade	8
One-shot: <code>asyncio.run()</code>	9
Repeated calls: <code>anyio.from_thread.start_blocking_portal()</code>	9
Where to go next	9
Profiles and auth	11
What "profile" means in kneo-client	11
Resolution order	11
TOML format	12
Environment variables	12
Auth schemes — what the platform accepts	13
Multi-profile workflows	13
Inspecting a resolved profile	14
Profile errors	15
Default config path	15
Idempotency and retries	17
What "idempotency and retries" mean in kneo-client	17
Idempotency keys, by default	17
When to supply your own key	18
Catching 409 mismatch	19
The retry policy	19
When retries fire	20
Customizing the policy	21
What surfaces when retries exhaust	22
Controlling the key (there is no opt-out)	22
Putting it together — a robust pattern	23
List-method results	24
Per-method return shapes	24
Page[T]	25
Iterating one page	25

Walking all pages: <code>iterate_all()</code>	26
window and the true total	26
Audit is fully paginated	27
Choosing a <code>page_size</code>	27
<code>Map[K, V]</code>	27
Filters + paging	28
Drop to the raw response	29
Auto-walking is not on the roadmap	29
Polling and waiting	30
Async create (<code>async_mode=true</code>)	
<code>wait_for_completion()</code>	30
<code>tail_trace()</code>	31
When to use which	32
Polling cadence	32
Custom terminal statuses	33
TimeoutError handling	33
Error handling	35
What "errors" mean in kneo-client	35
Hierarchy	35
What every exception carries	36
Status code → exception mapping	36
Catching patterns	39
Catch broadly, log richly	39
Branch on specific status	39
Branching on <code>KneoError.code</code>	40
Handle idempotency-key mismatches loudly	42
Don't catch successful-status branches	43
Transport errors and network troubleshooting	43
Timeouts	43
TLS verification	44
HTTP proxies	44
Inspecting the underlying <code>httpx</code> exception	45
Building error messages for users	45
Why a typed hierarchy	46
Reference	46
Compatibility matrix	48
What "compatibility" means in kneo-client	48
Current matrix	48
How pinning works in practice	49
Forward compatibility — newer <code>kneo_serv</code> than the pin	49

Backward compatibility — older kneo_serv than the pin 50

Server-version floors and behavior changes 51

Breaking changes 52

Verifying compatibility yourself 52

What the pin does not guarantee 53

Quickstart

Source: <docs/user/quickstart.md>

This guide walks you from a fresh shell to your first authenticated call against a running Kneo Agent Platform instance, then to creating a run and inspecting its trace.

What "quickstart" means in kneo-client

`kneo-client` is a typed Python SDK and adapter toolkit for the Kneo Agent Platform's `/v1` HTTP API. Two products use it as their shared client layer — Kneo Agent Dashboard for operations and Kneo Agent Studio for development — but you can use it directly from any Python 3.12+ codebase that needs to talk to a platform instance.

A working setup needs three things:

1. The `kneo-client` package installed.
2. A **profile** — a `(url, api_key, auth_scheme, timeout)` tuple resolved from a TOML config file, environment variables, or explicit kwargs. See [Profiles and auth](#) for the full resolution model.
3. A reachable platform instance.

The rest of this page assumes you have all three.

Install

```
python -m pip install kneo-client
```

Requires Python ≥ 3.12 (3.12 / 3.13 / 3.14 tested in CI). The runtime closure is small: `httpx`, `pydantic`, `anyio`, `platformdirs`, `attrs`. No CLI is installed — `kneo-client` is a library, not a service.

Verify the install:

```
python -c "from kneo_client import KneoClient, __version__; print(__version__)"
# → 0.8.0
```

Configure a profile

The client reads connection details in this order (later sources override earlier):

1. A TOML config file at `~/.config/kneo/client.toml` (XDG-style location via [platformdirs](#)).
2. Environment variables: `KNEO_PROFILE`, `KNEO_URL`, `KNEO_API_KEY`, `KNEO_AUTH_SCHEME`, `KNEO_TIMEOUT`.
3. Explicit keyword arguments to `KneoClient.from_profile()` (or the underlying `load_profile()`).

The minimum-friction setup is env vars — useful for one-off scripts and CI:

```
export KNEO_URL=https://kneo.example.com
export KNEO_API_KEY=your-api-key
```

For named multi-profile setups (e.g. `default` for prod, `staging` for non-prod), use the TOML file:

```
# ~/.config/kneo/client.toml
[default]
url = "https://kneo.example.com"
api_key = "prod-key"
auth_scheme = "bearer"
timeout = 30.0

[staging]
url = "https://staging-kneo.example.com"
api_key = "staging-key"
```

Switch profiles by name (`KneoClient.from_profile("staging")`) or by env (`KNEO_PROFILE=staging`).

See [Profiles and auth](#) for the full resolution semantics, both auth schemes, and the multi-profile workflow.

Ping the platform

The smallest possible script — verifies that auth flows correctly and the platform is reachable:

```
import asyncio
from kneo_client import KneoClient

async def main():
    async with KneoClient.from_profile() as client:
        ready = await client.platform.health.readyz()
        print(f"platform ready: ok={ready.ok} version={ready.service!r}")
```

```
asyncio.run(main())
```

What this exercises:

- Profile resolution (TOML / env / kwargs).
- API-key injection on the outgoing request (`Authorization: Bearer ...` or `X-Kneo-API-Key: ...` depending on scheme).
- The transport's retry loop (will retry on transient network failures).
- Error mapping (a missing key surfaces as `KneoAuthError`, an unreachable URL as `KneoNetworkError`).

If you get a `KneoAuthError`, your API key isn't reaching the platform — re-check the profile resolution chain. If you get a `KneoNetworkError`, the URL is unreachable; see [Error handling § Transport errors and network troubleshooting](#) for timeouts, TLS, proxies, and how to inspect the underlying httpx cause.

Create your first run

A run is the unit of work the platform executes — it instantiates a spec (an agent definition) and tracks its lifecycle through `queued` → `running` → `terminal` {`completed`, `failed`, `cancelled`, `timed_out`, `expired`}. To create one:

```
async with KneoClient.from_profile() as client:
    created = await client.platform.runs.create({"spec_id": "your-spec-id"})
    print(f"run_id={created.run_id} status={created.status}")

    terminal = await client.platform.runs.wait_for_completion(
        created.run_id, poll_interval=2.0, timeout=600
    )
    print(f"final status: {terminal.status}")

    trace = await client.platform.runs.trace(created.run_id, limit=20)
    for event in trace: # trace is a Page – iterate it directly (or use trace.items)
        print(event)
```

The interesting parts:

- `runs.create(...)` auto-injects an `Idempotency-Key` header (UUID4). Re-running the same payload with the same key is safe — the platform replays the original response. See [Idempotency and retries](#).
- `runs.wait_for_completion(...)` polls `runs.get(run_id)` until the run reaches a terminal status. Default terminal set is `{"completed", "failed", "cancelled", "timed_out", "expired"}` — the platform's canonical terminal set; pass `terminal_statuses={"blocked", ...}` to also stop when a run pauses for human review.

- `runs.trace(...)` returns events the platform recorded during the run (tool calls, model calls, middleware decisions, policy checks, etc.).

This is the operational core of the platform adapter. Every other endpoint follows the same shape: `client.platform.<resource>.<method>(...)` returning a typed model.

Skills catalog and per-request overlays

The agent surface (`client.agent.*`) mirrors the same shape. One pairing worth knowing early: the platform publishes a **skill catalog** — the declared + default skills a spec may reference by name — and `runs.create` accepts a per-request **skills overlay** that adds or disables skills for just that run:

```
async with KneoClient.from_profile() as client:
    catalog = await client.agent.skills.list() # Page; requires kneo_serv >= 0.8.0
    for skill in catalog:
        print(skill["name"], skill["description"])

    created = await client.platform.runs.create({
        "spec_id": "your-spec-id",
        "skills": {"add": ["web_search"], "disable": ["code_exec"]},
    })
```

The `skills` overlay on `runs.create` requires `kneo_serv >= 0.9.0`, and every name in `add` must be declared in the spec's skills block — the server rejects undeclared names. Use `agent.skills.list()` to discover the valid targets.

Sync facade

If your caller can't run an event loop (a script, a notebook, a sync framework), wrap the transport with `SyncTransport`:

```
from kneo_client.core.profiles import load_profile
from kneo_client.core.transport import SyncTransport

with SyncTransport(load_profile()) as transport:
    response = transport.request("GET", "/v1/healthz")
    print(response.json())
```

Under the hood `SyncTransport` runs `Transport` inside an `anyio.from_thread.start_blocking_portal()` — same retry / idempotency / error flows, called synchronously. The async surface is the recommended path; the sync facade is for ergonomics.

`SyncTransport` does **not** mount `.platform / .agent` adapters — those are async-only. Two patterns work for sync consumers who want the wrapped endpoints:

One-shot: `asyncio.run()`

If you're calling once and not in an existing event loop:

```
import asyncio
from kneo_client import KneoClient

async def fetch() -> str:
    async with KneoClient.from_profile() as client:
        run = await client.platform.runs.create({"spec_id": "s1"})
        return run.run_id

run_id = asyncio.run(fetch())
print(run_id)
```

Repeated calls: `anyio.from_thread.start_blocking_portal()`

If you're calling many times from sync code (e.g., inside a synchronous CLI loop) and want to avoid the startup cost of `asyncio.run()` per call, hold open a portal — the same machinery `SyncTransport` itself uses:

```
from anyio.from_thread import start_blocking_portal
from kneo_client import KneoClient

with start_blocking_portal() as portal:
    client = portal.call(KneoClient.from_profile().__aenter__)
    try:
        run = portal.call(client.platform.runs.create, {"spec_id": "s1"})
        status = portal.call(client.platform.runs.get, run.run_id)
        print(status.status)
    finally:
        portal.call(client.__aexit__, None, None, None)
```

The portal runs an event loop on a background thread; `portal.call(async_fn, *args)` blocks the calling thread until the coroutine completes. This is the recommended pattern for long-lived sync code that talks to the platform repeatedly.

Where to go next

By topic:

- [Profiles and auth](#) — multi-profile workflows, the two header schemes the platform supports, environment-variable precedence.
- [Idempotency and retries](#) — when keys are auto-injected, how 409 mismatches surface, customizing the retry policy.
- [List-method results](#) — `Page[T]` and `Map[K, V]` wrappers, walking pages with `iterate_all()` (audit included — fully paginated since `kneo_serv` 0.6.0), the paging `window / true-total` semantics.
- [Error handling](#) — the full exception hierarchy and what to catch.
- [Compatibility matrix](#) — which `kneo-client` releases support which `kneo_serv` versions.

For the comprehensive API reference (every class, method, exception), see the [API Reference HTML](#) or the [PDF version](#).

For runnable end-to-end scripts: [examples/](#) .

Profiles and auth

Source: docs/user/profiles_and_auth.md

This guide explains how `kneo-client` resolves connection details (URL, API key, auth scheme, timeout) into a profile, what the two API-key header schemes mean, and how to set up multi-profile workflows for dev / staging / prod.

What "profile" means in kneo-client

A **profile** is a frozen dataclass bundling `(name, url, api_key, auth_scheme, timeout)`:

```
@dataclass(frozen=True)
class Profile:
    name: str
    url: str
    api_key: str
    auth_scheme: AuthScheme = AuthScheme.BEARER
    timeout: float = 30.0
```

Everything `Transport` needs to talk to one Kneo Agent Platform instance fits in those five fields. A `KneoClient` is built around exactly one profile; you construct one per platform instance you talk to.

Profiles are typically named — `default`, `staging`, `prod`, `ci`, etc. The name is informational (it appears in log lines) but doubles as the section name in the TOML config file.

Resolution order

`load_profile()` and `KneoClient.from_profile()` merge values from three sources, with later sources overriding earlier ones:

1. **TOML config file**, by default `~/.config/kneo/client.toml` (XDG-style via [platformdirs](#)). Pass `config_file=Path(...)` to point at a different file. If the file doesn't exist, it's skipped silently — env vars and explicit kwargs are still consulted.
2. **Environment variables**: `KNEO_URL`, `KNEO_API_KEY`, `KNEO_AUTH_SCHEME`, `KNEO_TIMEOUT`. (`KNEO_PROFILE` selects which TOML section to load — it does *not* override field values.)
3. **Explicit keyword arguments** to the function call.

If `url` or `api_key` cannot be resolved from any source, a `ProfileError` is raised with details on which sources were checked. Bad TOML, an unknown `auth_scheme`, or a non-numeric `timeout` also surface as `ProfileError`.

```

from kneo_client.core.profiles import load_profile

p = load_profile()                # 'default' from TOML + env
p = load_profile("staging")       # explicit profile name
p = load_profile(url="https://ad-hoc", api_key=token) # explicit kwargs win

```

TOML format

Each top-level table is one profile:

```

# ~/.config/kneo/client.toml
[default]
url = "https://kneo.example.com"
api_key = "prod-key"
auth_scheme = "bearer"          # or "kneo_api_key"
timeout = 30.0                  # seconds

[staging]
url = "https://staging-kneo.example.com"
api_key = "staging-key"

[local]
url = "http://127.0.0.1:8000"
api_key = "dev-token"
auth_scheme = "kneo_api_key"

```

`auth_scheme` and `timeout` are optional; their defaults are `"bearer"` and `30.0`.

Picking a profile at call time:

```

client = KneoClient.from_profile()      # 'default' (or $KNEO_PROFILE)
client = KneoClient.from_profile("staging") # explicit
client = KneoClient.from_profile("local")  # explicit

```

Environment variables

Variable	Purpose
<code>KNEO_PROFILE</code>	Profile name to load. Falls back to <code>"default"</code> if unset.
<code>KNEO_URL</code>	Override the profile's URL.

Variable	Purpose
<code>KNEO_API_KEY</code>	Override the profile's API key.
<code>KNEO_AUTH_SCHEME</code>	Override the scheme. Accepts <code>"bearer"</code> or <code>"kneo_api_key"</code> .
<code>KNEO_TIMEOUT</code>	Override the per-request timeout (float seconds).

CI environments typically set just `KNEO_URL` and `KNEO_API_KEY` and skip the TOML file entirely. The lack of a config file is *not* an error — env vars + kwargs can satisfy resolution on their own.

A bad value for `KNEO_TIMEOUT` (non-numeric) raises `ProfileError` with the variable name in the message — easier to debug than a silent fallback.

Auth schemes — what the platform accepts

The platform accepts the API key in either of two header schemes. They are *semantically* equivalent — both end up at the same platform code path — but operationally they have different trade-offs:

Scheme	Header sent	When to choose
<code>bearer</code> (default)	<code>Authorization: Bearer <key></code>	Works with most reverse proxies. Easy to revoke at the gateway layer. Standard HTTP semantics.
<code>kneo_api_key</code>	<code>X-Kneo-API-Key: <key></code>	Useful when your edge stack already uses the <code>Authorization</code> header for something else (mutual TLS auth, an upstream OAuth flow, etc.). Avoids the collision.

When in doubt, start with `bearer`. You can switch schemes per-profile without code changes — just update the TOML or the env var.

The two schemes are implemented by `kneo_client.core.auth.ApiKeyAuth`, an `httpx.Auth` subclass that injects whichever header the active profile selects. Internally the auth flow runs *inside* `httpx`'s request flow (after `Transport` has added its other headers), so the API key reaches every redirect / retry attempt at the right layer.

Multi-profile workflows

A common pattern in CI / local dev:

```
import os
from kneo_client import KneoClient

profile_name = "ci" if os.getenv("CI") else "default"
async with KneoClient.from_profile(profile_name) as client:
    ...
```

Or override explicitly when the situation calls for an ad-hoc connection:

```
async with KneoClient.from_profile(url="https://ad-hoc.example.com", api_key=tok) as client:
    ...
```

Explicit kwargs always win, so you can keep the TOML file as a baseline and override per-call.

Programmatic profile construction (when secrets come from a vault / secrets manager) skips `load_profile()` entirely:

```
from kneo_client import KneoClient
from kneo_client.core.auth import AuthScheme
from kneo_client.core.profiles import Profile

def profile_from_vault() -> Profile:
    secret = vault.get("kneo/prod")
    return Profile(
        name="prod",
        url=secret["url"],
        api_key=secret["api_key"],
        auth_scheme=AuthScheme.BEARER,
        timeout=30.0,
    )

async with KneoClient(profile_from_vault()) as client:
    ...
```

`Profile` is a frozen dataclass — pass it directly to `KneoClient(profile)`.

Inspecting a resolved profile

`KneoClient.profile` returns the `Profile` actually in use:

```
async with KneoClient.from_profile() as client:
    print(f"connected to {client.profile.url} as profile {client.profile.name!r}")
```

The `api_key` field is on the dataclass — handle it like any other secret. The redaction-aware logger in `kneo_client.core.logging` masks the key whenever it logs request / response headers. The dataclass itself is also redacted in its `repr`: `api_key` is declared `field(repr=False)`, so `repr(profile)` / printing the profile does **not** expose the key. You still shouldn't pass `profile.api_key` to a log sink yourself — `repr=False` only guards the default representation, not explicit access.

If you want to log a profile's non-secret fields explicitly:

```
print(f"profile name={p.name!r} url={p.url!r} scheme={p.auth_scheme.value!r} timeout={p.timeout}")
```

...and never include `p.api_key` in anything that goes to a log sink.

Profile errors

`ProfileError` covers five failure modes:

Trigger	Message pattern
Missing <code>url</code> after all sources	"profile 'X': 'url' is not set ..."
Missing <code>api_key</code> after all sources	"profile 'X': 'api_key' is not set ..."
Malformed TOML	"failed to parse <path>: ..."
Unknown auth scheme	"unknown auth_scheme '...'; expected one of: bearer, kneo_api_key"
Non-numeric <code>KNEO_TIMEOUT</code>	"\$KNEO_TIMEOUT must be a float, got '...'"

All five are explicit and name the offending source. Catch `ProfileError` (or `Exception` if you don't care which) at process startup to fail fast with a clear message rather than blowing up on the first call.

Default config path

`kneo_client.core.profiles.default_config_path()` returns the XDG-style default — `~/.config/kneo/client.toml` on Linux, `~/Library/Application Support/kneo/client.toml` on macOS, the appropriate `%APPDATA%\kneo\client.toml` on Windows. Resolved via `platformdirs.user_config_dir("kneo")`.

If you want a project-local TOML (committed to a repo, picked up by CI without needing a user-config), pass it explicitly:

```
from pathlib import Path  
client = KneoClient.from_profile(config_file=Path(".kneo.toml"))
```


Idempotency and retries

Source: docs/user/idempotency_and_retries.md

This guide explains when `kneo-client` injects idempotency keys, how 409 mismatches surface, what the retry policy is, and how to customize either piece for non-default deployments.

What "idempotency and retries" mean in kneo-client

The Kneo Agent Platform's write endpoints are designed for *safe retry* where the spec documents the `Idempotency-Key` header: run create and continue, `specs/run`, and human-task resume (`continue` is contractual on `kneo_serv >= 0.9.0`; older servers treated its key as best-effort). The platform short-circuits a duplicate `POST` with the same `Idempotency-Key` header and identical payload — the second request replays the original response rather than re-executing the side effect. One exception worth knowing: `runs.cancel` has *no* documented key (even on `kneo_serv 0.9.0`) — a retried cancel may be observed twice, which is harmless because cancel is naturally idempotent in effect.

`kneo-client` makes that safety automatic by:

1. Auto-injecting a fresh UUID4 `Idempotency-Key` on every `POST` (unless the caller supplies their own).
2. Retrying transient transport errors and the fixed set `{429, 502, 503, 504}` within a configurable `RetryPolicy`.
3. Honoring `Retry-After` on 429 and 503 responses (the server's hint overrides the policy's computed delay; a hint above `RetryPolicy.max_delay` fails fast instead — see below). 503 is `kneo_serv`'s run-queue backpressure signal on `POST /v1/runs`; because `POSTs` carry an auto-injected idempotency key, those are retry-safe and retried automatically.
4. Surfacing the platform's payload-mismatch behavior as `KneoIdempotencyMismatchError` so retry-with-wrong-payload bugs are loud, not silent.

Both pieces work together: idempotency keys make it *safe* to retry a `POST`; the retry policy *decides* when to retry. The default settings work for most callers; you can override either independently.

Idempotency keys, by default

On every `POST` the transport sends, it adds:

- `Idempotency-Key`: `<fresh UUID4>` — unless the caller passes `idempotency_key=<string>` on the method call.

The platform's contract:

- **Same key + identical payload** → server short-circuits and returns the original response. The retry is effectively a replay.
- **Same key + original attempt still executing** → server returns HTTP 409 with code `idempotency_key_in_progress` and a `Retry-After` hint. The transport auto-retries this case per the hint; it surfaces as a `KneoConflictError` only once retries exhaust.
- **Same key + different payload** → server returns HTTP 409 with code `idempotency_key_conflict`. The client surfaces this as `KneoIdempotencyMismatchError`.
- **Different key** → server treats this as a new request and executes the side effect.

The auto-generated key is a fresh UUID4 per request. Collisions are not a real concern (UUID4 has 122 bits of randomness), so by default each call is independent — retries beyond the transport's own loop won't be deduplicated by the server.

When to supply your own key

The auto-generated key is fine for one-shot calls. Supply your own when:

- **You're retrying outside the transport's loop.** The transport retries within `RetryPolicy.max_attempts` for transient failures, but if your *application* catches an error and retries (e.g. a job runner re-invoking after a process restart), pass the same key on both attempts so the platform dedupes. The transport's retries already share the key.
- **You want cross-process correlation.** Two services submitting the same logical request can dedupe by agreeing on the key (e.g. derive it from a hash of the request payload + a request ID from your application).
- **You're testing.** A stable key makes test fixtures deterministic.

```
from kneo_client.core.idempotency import new_idempotency_key

key = new_idempotency_key()
body = {"spec_id": "my-spec"}

# First attempt – succeeds normally
run = await client.platform.runs.create(body, idempotency_key=key)

# ... later, retrying after a transient outage outside the transport's loop:
run = await client.platform.runs.create(body, idempotency_key=key)
# → returns the first run (same response, no new side effect)
```

Constraints on caller-supplied keys (validated + normalized by the client before sending):

- **Non-empty, non-whitespace** — an empty or whitespace-only key raises `ValueError`. The platform strips surrounding whitespace and treats a whitespace-only key as *no* key — which

would silently disable idempotency (a retry re-runs the side effect), so the client rejects it up front.

- **Stripped before send** — surrounding whitespace is removed, so a padded key `" abc "` is transmitted as `"abc"`. This matches the platform's stored key, so the value you see echoed back (including `KneoError.idempotency_key`) is the stripped form.
- **At most 256 characters** (`MAX_KEY_LENGTH`), measured on the *stripped* value. The platform enforces this (rejecting an over-length key with `400 invalid_idempotency_key`); the client validates locally to fail faster.

Catching 409 mismatch

The platform's 409s carry a stable envelope code (`KneoError.code`), and the client classifies on it. Only `idempotency_key_conflict` — the *same key reused with a different payload* — raises `KneoIdempotencyMismatchError` (plus, against pre-0.6.0 servers whose bodies carry no code, any 409 on a keyed request). Other codes raise the `KneoConflictError` base: `run_state_conflict` (a lifecycle fence, e.g. cancelling an already-terminal run) and `resource_locked` (held by another operation), both carrying `.retry_after` when the server hints one. `idempotency_key_in_progress` is auto-retried by the transport per its `Retry-After` and only surfaces as `KneoConflictError` once retries exhaust.

A mismatch is almost always a caller bug. Surface it loudly:

```
from kneo_client.core.errors import KneoIdempotencyMismatchError

try:
    await client.platform.runs.create(payload, idempotency_key=key)
except KneoIdempotencyMismatchError as exc:
    raise RuntimeError(
        f"Idempotency-Key {exc.idempotency_key!r} was reused with a different payload. "
        f"Either generate a new key for the new request or fix the payload drift."
    ) from exc
```

The exception carries the key as sent (`exc.idempotency_key`), the platform's error body (`exc.body`), the server-assigned request ID for log correlation (`exc.request_id`), and the HTTP status (`exc.status == 409`).

Note that `KneoIdempotencyMismatchError` is a *subclass* of `KneoConflictError`, so a generic `except KneoConflictError` will also catch it. If you only care about non-mismatch conflicts (`run_state_conflict`, `resource_locked`), catch `KneoIdempotencyMismatchError` first — or branch on `exc.code`.

The retry policy

`RetryPolicy` is a frozen dataclass describing when and how long to wait between attempts. The transport applies it; the policy itself does no I/O.

The default policy:

```
RetryPolicy(
    max_attempts=3,    # up to 3 total attempts
    base_delay=0.2,    # ~0.2s before attempt 2
    max_delay=30.0,    # cap on the computed delay
    jitter=0.1,        # 10% jitter applied to the computed delay
)
```

Delay sequence (without jitter): attempt 1 → no delay, attempt 2 → `base_delay`, attempt 3 → `2 × base_delay`, ..., capped at `max_delay`. With `jitter=0.1`, a 1-second delay becomes uniformly distributed in `[0.9, 1.1]`.

`Retry-After` from a 429 or 503 response overrides the computed delay — no jitter is applied on top — when the hint is within `max_delay`. A hint **above `max_delay` is not slept on at all**: retrying earlier than the server asked is likely doomed, so the transport stops retrying immediately and raises the typed error with the verbatim hint on `.retry_after` — an intermediary advertising an hour-long `Retry-After` neither blocks the caller nor burns attempts on guaranteed-fail retries. Apply your own longer back-off from `.retry_after` when you want to wait it out.

When retries fire

The transport retries on:

- Transport-level errors from `httpx` — DNS resolution failures, connect failures, TLS handshakes, read timeouts.
- HTTP **429** (rate limited), **502** (bad gateway), **503** (service unavailable), **504** (gateway timeout). Other 4xx and 5xx status codes do *not* trigger a retry — those typically indicate caller errors or non-transient server problems that retrying won't fix.
- HTTP **409** with code `idempotency_key_in_progress` — the server saying the key's original attempt is still executing; retried per its `Retry-After`. Every other 409 is a real conflict and surfaces immediately.

The retryable status set is `RETRYABLE_STATUS_CODES = frozenset({429, 502, 503, 504})` — a module constant, intentionally not configurable. If your platform deployment legitimately returns transient 500 s, fix the deployment; we'd rather diagnose the root cause than open the gate.

Status-code retries fire only for:

- **Idempotent verbs**: `GET`, `HEAD`, `OPTIONS`, `PUT`, `DELETE`.

- **POST with an Idempotency-Key** — i.e. always, since the transport auto-injects one on every POST.

Transport-level errors are slightly stricter: they're retried only for idempotent verbs, keyed POSTs, or — for any method — *connect-phase* errors (`httpx.ConnectError` and `httpx.ConnectTimeout` : connection refused or the TCP/TLS handshake timing out), which provably never reached the server and are therefore always safe to re-send. Mid-flight decode failures (`httpx.DecodingError` , e.g. a proxy emitting corrupt gzip) and redirect loops (`httpx.TooManyRedirects`) also count as transient — but only on replay-safe methods, since the request may have executed server-side.

This is why auto-injection matters: it's what makes POST retries safe.

Customizing the policy

Pass a custom policy when constructing the client:

```
from kneo_client import KneoClient
from kneo_client.core.profiles import load_profile
from kneo_client.core.retries import RetryPolicy

profile = load_profile()

# Aggressive retries for a flaky network
policy = RetryPolicy(max_attempts=8, base_delay=0.5, max_delay=60.0)
client = KneoClient(profile, retry_policy=policy)

# Flat delay (no exponential growth)
policy = RetryPolicy(max_attempts=3, base_delay=2.0, max_delay=2.0, jitter=0)
client = KneoClient(profile, retry_policy=policy)
```

To disable retries entirely (useful in tests that want to see the first failure surface immediately):

```
client = KneoClient(profile, retry_policy=RetryPolicy(max_attempts=1))
```

Constraints on `RetryPolicy` parameters (validated at construction):

- `max_attempts ≥ 1`
- `base_delay ≥ 0`
- `max_delay ≥ base_delay`
- `0 ≤ jitter ≤ 1`

Invalid values raise `ValueError` immediately — you find out at startup, not on the first retry.

What surfaces when retries exhaust

If all retries fail, the client raises the appropriate typed exception based on the *final* attempt's outcome:

Final attempt failed because...	Exception raised
HTTP 429	<code>KneoRateLimited</code> (carries <code>.retry_after</code>)
HTTP 503	<code>KneoServiceUnavailable</code> (a <code>KneoServerError</code> subclass; carries <code>.retry_after</code>)
HTTP 502 / 504	<code>KneoServerError</code>
Transport error (DNS, connect, TLS, read)	<code>KneoNetworkError</code>

Intermediate retry attempts are logged at `INFO` level under the `kneo_client.transport` logger:

```
INFO:kneo_client.transport:status 503 on attempt 1; sleeping 0.20s
INFO:kneo_client.transport:status 503 on attempt 2; sleeping 0.40s
```

If you want to see the retry behavior in your application logs, set the logger level:

```
import logging
logging.basicConfig(level=logging.INFO)
logging.getLogger("kneo_client.transport").setLevel(logging.INFO)
```

Controlling the key (there is no opt-out)

The transport auto-injects an `Idempotency-Key` on every `POST` *unconditionally* — and it overwrites any `Idempotency-Key` you place in a `headers=` mapping on `Transport.request`, so passing explicit headers is **not** an escape hatch. The supported way to control the key is the `idempotency_key=` kwarg on the wrapper methods (or on `Transport.request` directly):

```
# Standard call – transport adds a fresh UUID4 automatically
await client.platform.runs.create(body)

# Deterministic key – full control over what gets sent
await client.platform.runs.create(body, idempotency_key="your-deterministic-key")
```

There is no supported way to send a `POST` *without* a key. The platform's contract assumes the header is present, and the auto-injection is exactly what makes `POST` retries safe — an opt-out would quietly disable that.

Putting it together — a robust pattern

```

from kneo_client import KneoClient
from kneo_client.core.errors import (
    KneoIdempotencyMismatchError, KneoNetworkError, KneoServerError,
)
from kneo_client.core.idempotency import new_idempotency_key

async def create_run_robust(client: KneoClient, body: dict) -> str:
    key = new_idempotency_key() # one key, used across application-level retries

    for attempt in range(1, 4):
        try:
            run = await client.platform.runs.create(body, idempotency_key=key)
            return run.run_id

        except KneoIdempotencyMismatchError as exc:
            # Caller bug – body changed between attempts. Don't retry.
            raise RuntimeError(f"payload drift on key {exc.idempotency_key!r}") from exc

        except (KneoNetworkError, KneoServerError) as exc:
            # The transport already retried within its policy; we add an outer
            # guard for application-level recovery (e.g. across process restarts).
            if attempt == 3:
                raise
            await asyncio.sleep(2 ** attempt)

    raise AssertionError("unreachable")

```

In practice the transport's built-in retries are sufficient for most use cases; the outer loop above is for the rare case where you want application-level retry behavior (e.g. survive a process restart while the platform is temporarily unreachable).

List-method results

Source: docs/user/list_results.md

Every `.list()` -style method on `kneo_client` returns a typed wrapper around the server response. There are two wrapper classes — chosen by the underlying server-response shape, not by "does this paginate":

- **Page[T]** for *list-shaped* responses (runs, checkpoints, trace events, human tasks, audit events, the skill catalog).
- **Map[K, V]** for *dict-keyed* responses (credentials inventory, per-environment policies).

Both live in `kneo_client.core.results` and are exported from `kneo_client.core`. The free-function `iterate_all()` async iterator is also there and walks pages for `Page` returns whose endpoint supports it.

Per-method return shapes

Method	Returns	Server-side pagination
<code>client.platform.runs.list()</code>	<code>Page[Run]</code>	Full <code>limit / offset / total</code>
<code>client.platform.runs.checkpoints(run_id)</code>	<code>Page[Checkpoint]</code>	Full <code>limit / offset / total</code>
<code>client.platform.runs.trace(run_id)</code>	<code>Page[TraceEvent]</code>	Full <code>limit / offset / total</code>
<code>client.platform.human_tasks.list()</code>	<code>Page[HumanTask]</code>	Full <code>limit / offset / total</code>
<code>client.platform.audit.list()</code>	<code>Page[AuditEvent]</code>	Full <code>limit / offset / total</code> (since <code>kneo_serv</code> 0.6.0; see below)
<code>client.agent.skills.list()</code>	<code>Page[Skill]</code>	Full <code>limit / offset / total</code> (endpoint requires <code>kneo_serv</code> $\geq 0.8.0$)
<code>client.platform.credentials.list()</code>	<code>Map[str, Credential]</code>	None — whole map returned in one call

Method	Returns	Server-side pagination
<code>client.platform.policies.environment_list()</code>	<code>Map[str, EnvironmentPolicy]</code>	None — whole map returned in one call

The item / value types (`Run`, `Checkpoint`, `Credential`, etc.) are the generated models from `kneo_client._generated.models`. Several are *open* objects in the pinned spec (no declared fields), in which case fields are reached via bracket lookup — `run["run_id"]` rather than `run.run_id`. The [API Reference](#) lists which per-endpoint item models are typed vs open.

Page[T]

```
@dataclass(frozen=True)
class Page(Generic[T]):
    items: list[T]
    total: int | None = None
    limit: int | None = None
    offset: int | None = None
    sort_by: str | None = None
    sort_order: str | None = None
    window: int | None = None

    @property
    def count(self) -> int: ...          # = len(items)

    @property
    def has_more(self) -> bool: ...      # False when offset/total are None; clamps to window

    def __iter__(self) -> Iterator[T]: ...
    def __len__(self) -> int: ...
    def __getitem__(self, index: int) -> T: ...
```

Metadata fields are Optional — a server that doesn't echo a particular field (e.g. a pre-0.6.0 `kneo_serv` answering `audit.list()`, or a pre-0.9.0 one for `window`) leaves it as `None` rather than carrying a synthesized value. `has_more` derives sensibly: `False` whenever `offset` or `total` is `None`, and clamped to `window` when the server discloses one (see [below](#)).

Iterating one page

`Page` is a sequence; iterate it directly:

```
page = await client.platform.runs.list(status="running", limit=50)
print(f"got {page.count} of {page.total} runs")
```

```
for run in page:
    print(run["run_id"])
```

Indexing and `len` work as expected:

```
first_run = page[0]
total_on_this_page = len(page)
```

Walking all pages: `iterate_all()`

For the fully-paginating endpoints, `iterate_all()` walks pages transparently. Pass a `fetch_page(limit, offset)` callable that returns the next `Page`:

```
from kneo_client.core.results import iterate_all

async def fetch(limit: int, offset: int):
    return await client.platform.runs.list(limit=limit, offset=offset, status="running")

async for run in iterate_all(fetch, page_size=200):
    print(run["run_id"])
```

`iterate_all()` does three things on your behalf:

1. **Clamps `page_size` to `MAX_PAGE_SIZE = 1000`** (the platform's hard upper bound). Asking for more silently downsizes.
2. **Walks `offset` automatically** — each page's `offset` plus its `count` becomes the next page's starting position.
3. **Stops when `has_more` is `False`** — either because the server has no more *fetchable* data (`offset + count ≥ total`, clamped to the paging `window` — see next section) or because the server doesn't echo enough metadata to know (pre-0.6.0 audit, see below).

`window` and the true `total`

Since `kneo_serv` 0.9.0, list responses disclose `window` — the deepest paging offset the server will reach — and `total` is the true store-side count (`COUNT(*)`), which can *exceed* it. When both are present, `has_more` reports whether more items are **fetchable**, not merely whether they exist: it clamps `total` to `window`, so pagination loops (including `iterate_all()`) stop honestly at the window edge instead of fetching a guaranteed-empty page. To detect items beyond the reachable window, compare `page.total` against `page.window`:

```
page = await client.platform.runs.list(limit=100)
if page.window is not None and page.total is not None and page.total > page.window:
    print(f"{page.total - page.window} runs exist beyond the paging window")
```

Pre-0.9.0 servers don't echo `window`; it stays `None` and `has_more` falls back to the plain `offset + count < total` derivation.

Audit is fully paginated

`audit.list()` has been fully paginated since `kneo_serv 0.6.0 / kneo-client 0.6.0`: the endpoint accepts `limit / offset / sort_by / sort_order`, the response echoes `total / offset` and sort metadata, and `iterate_all()` walks it like any other endpoint:

```
page = await client.platform.audit.list(event_type="run.created", limit=200)
print(f"got {page.count} of {page.total} matching audit events")

async def fetch(limit: int, offset: int):
    return await client.platform.audit.list(
        event_type="run.created", limit=limit, offset=offset
    )

async for event in iterate_all(fetch, page_size=200):
    print(event)
```

One back-compat note: pre-0.6.0 servers omit the audit pagination metadata, in which case the fields degrade to `None`, `has_more` is `False`, and `iterate_all()` fetches once and exits.

Choosing a `page_size`

Rules of thumb for the fully-paginating endpoints:

- **100–200** for most interactive flows — snappy first-byte latency.
- **500–1000** for back-end exports — fewer requests, higher per-request cost on retry.
- **< 50** when per-item downstream processing is slow and you want to start producing output sooner.

The hard ceiling is `MAX_PAGE_SIZE = 1000`; `iterate_all()` clamps for you.

Map[K, V]

```
@dataclass(frozen=True)
class Map(Generic[K, V]):
    items: Mapping[K, V]

    @property
    def count(self) -> int: ...          # = len(items)

    def __getitem__(self, key: K) -> V: ...
```

```

def __iter__(self) -> Iterator[V]: ...    # iterates values
def __len__(self) -> int: ...
def __contains__(self, key: object) -> bool: ...

def get(self, key: K, default: V | None = None) -> V | None: ...
def keys(self) -> KeysView[K]: ...
def values(self) -> ValuesView[V]: ...

```

`Map` covers the two endpoints whose responses are keyed maps rather than item lists. The whole collection comes back in one call — there is no `limit` / `offset` to think about.

```

inv = await client.platform.credentials.list()
print(f"{inv.count} credential references")
for credential in inv:                    # iterates values
    print(credential["provider"])

aws = inv["cred-aws"]                    # bracket lookup
if "cred-missing" in inv:                # membership test
    ...

for cred_id in inv.keys():                # explicit keys view
    ...

# Pair iteration uses the underlying Mapping when needed:
for cred_id, body in inv.items.items():
    print(cred_id, body["provider"])

```

Iteration yields *values*, not keys — consistent with `Page` iterating items. Use `.keys()` or `.items.items()` if you need keys.

Filters + paging

Resource-specific filter kwargs compose with paging kwargs:

```

# Just the failures, walked in pages of 200
async def fetch(limit, offset):
    return await client.platform.runs.list(
        status="failed", limit=limit, offset=offset
    )

async for run in iterate_all(fetch, page_size=200):
    ...

# Audit events for one run, first 500 (paginate onward with offset=)
events = await client.platform.audit.list(run_id="r1", limit=500)

```

Drop to the raw response

If you need a field the wrapper doesn't expose — for example the `run_id` echo on `CheckpointListResponse` / `TraceResponse`, or a field the platform adds in a newer release — bypass the wrapper and call the transport directly:

```
from kneo_client._generated.models.checkpoint_list_response import (
    CheckpointListResponse,
)

resp = await client._transport.request(
    "GET", f"/v1/runs/{run_id}/checkpoints", params={"limit": 100}
)
raw = CheckpointListResponse.from_dict(resp.json())
print(raw.run_id, raw.checkpoints, raw.total)
```

`_transport` is the underlying `Transport`; the leading underscore signals it's an escape hatch, not part of the recommended surface. Anything you build on top of it tracks the generated layer directly and won't follow public-API compatibility — re-evaluate at each `kneo_serv` pin bump.

Auto-walking is *not* on the roadmap

`client.platform.runs.list_all()` returning an async iterator — i.e. auto-paged at the adapter layer — is intentionally *not* planned. The current shape (list methods return one `Page`; `iterate_all()` walks pages explicitly) keeps the per-call cost transparent and lets the caller decide when to stop. An accidental wide filter shouldn't walk millions of items behind the caller's back.

If you want a one-liner in your own code, write a small helper around `iterate_all()` and the relevant `.list()` callable — the pattern is identical for every fully-paginating endpoint.

Polling and waiting

Source: docs/user/polling_and_waiting.md

Two helpers on `RunsClient` cover the two common run-lifecycle waiting patterns:

- `wait_for_completion()` — block until a run reaches a terminal status, returning the final status. Best when you only care about the outcome and will inspect the trace afterwards.
- `tail_trace()` — stream trace events as they arrive, returning an async iterator. Best when you want to surface progress live (Studio editors, CLIs, dashboards).

Both follow the same polling-with-status-check pattern under the hood; they differ in what they yield to the caller.

Async create (`async_mode=true`)

A synchronous `runs.create()` blocks until the run finishes and returns its result. Passing `async_mode=true` instead tells the platform to accept the run and dispatch it in the background:

```
created = await client.platform.runs.create(
    {"spec_id": "my-spec", "async_mode": True}
)
status = await client.platform.runs.wait_for_completion(created.run_id)
```

Since `kneo_serv` 0.11.0 an async create responds with **HTTP 202 Accepted** (it was `200` through 0.10.x); a synchronous create still returns `200`. The client treats any `2xx` as success, so `create()` returns the same `RunCreateResponse` carrying the `run_id` in both cases — you don't branch on the status code. From there, poll with `wait_for_completion()` or stream with `tail_trace()` below. See examples/06_async_run.py for the end-to-end flow.

The agent-side `client.agent.specs.run()` (Studio's iterate-and-test route, `POST /v1/specs/run`) shares the same async semantics since `kneo_serv` 0.12.0/1.0.0: `{"async_mode": True}` returns a `202` with the queued `run_id`, which you poll with the same `wait_for_completion()`.

```
created = await client.agent.specs.run(
    {"spec_id": "my-spec", "async_mode": True}
)
status = await client.platform.runs.wait_for_completion(created.run_id)
```

`wait_for_completion()`

```
status = await client.platform.runs.wait_for_completion(
    run_id,
    poll_interval=1.0,  # seconds between GETs to /v1/runs/{run_id}
    timeout=600,        # total budget; None to wait indefinitely
)
print(status.status)    # "completed" | "failed" | "cancelled" | ...
```

The default `terminal_statuses` set is `{"completed", "failed", "cancelled", "timed_out", "expired"}` — the platform's canonical terminal set. `timed_out` (the platform's timeout sweep) and `expired` (a blocked human task's `on_timeout: fail` policy) became reachable in `kneo_serv` 0.9.0 and are terminal by default; no custom set is needed to stop on them. Pass a custom set to treat additional states as terminal:

```
# Stop as soon as the run pauses for human review.
status = await client.platform.runs.wait_for_completion(
    run_id,
    terminal_statuses={"blocked"},
)
```

Raises `TimeoutError` if the deadline elapses before any terminal status is reached.

tail_trace()

```
async for event in client.platform.runs.tail_trace(run_id, poll_interval=1.0):
    print(event["event_type"], event["payload"] if "payload" in event else None)
```

Trace-event items are generated "open" models, not dicts — they support bracket lookup (`event["key"]`) and membership tests (`"key" in event`), but not `dict` conveniences like `.get()`.

Internally `tail_trace` walks `/v1/runs/{run_id}/trace` with ascending offset and yields each event as it lands. When the run reaches a terminal status (configurable, same defaults as `wait_for_completion`) the helper does one final drain pass to capture events emitted between the last poll and the status transition, then returns.

Key arguments (all optional):

Arg	Default	Notes
<code>start_offset</code>	<code>0</code>	Resume tailing from a known position.

Arg	Default	Notes
<code>page_size</code>	<code>100</code>	<code>limit</code> per <code>/v1/runs/{run_id}/trace</code> fetch.
<code>poll_interval</code>	<code>1.0</code>	Seconds to sleep when no new events arrived.
<code>timeout</code>	<code>None</code>	Total budget for reaching terminal status. Does not apply during the post-terminal drain — that always runs to completion.
<code>terminal_statuses</code>	<code>{"completed", "failed", "cancelled", "timed_out", "expired"}</code>	Override to treat other states as terminal.
<code>event_type</code>	<code>None</code>	Filter passed through to <code>trace</code> .

Raises `TimeoutError` if the deadline elapses while waiting for a terminal status. Events yielded so far stay yielded — the caller has already consumed them.

When to use which

Want	Use
Just the final outcome	<code>wait_for_completion()</code>
Final outcome + the full trace at the end	<code>wait_for_completion()</code> then <code>runs.trace()</code>
Events as they arrive (live UI / log tail)	<code>tail_trace()</code>
First N events of a finished run	<code>runs.trace(limit=N)</code> directly
Resume an interrupted tail	<code>tail_trace(start_offset=...)</code>

Polling cadence

A rule of thumb: pick `poll_interval` so that *p99 perceived latency* roughly equals `poll_interval / 2`.

- **Interactive UIs (Studio, dashboards)** — 0.5 to 1.0 seconds. Snappy enough for human attention.
- **Operational scripts (CI, scheduled jobs)** — 2.0 to 5.0 seconds. Reduces request volume for long-running jobs where snappy feedback isn't needed.
- **Long-running background workers** — 10.0 + seconds. Pair with a generous timeout.

Don't go below 0.5 seconds without a specific reason — at that rate you're spending more on request overhead than you save in latency, and the platform's retry / backoff machinery is also factoring in.

Custom terminal statuses

The platform reports several non-final statuses (`running` , `queued` , `blocked` — paused for human review) that the default terminal set excludes.

The most common override is **treating human-review pauses as terminal** — useful when an operator-facing tool wants to surface "this run needs your attention" rather than continue polling:

```
status = await client.platform.runs.wait_for_completion(
    run_id,
    terminal_statuses={
        "completed", "failed", "cancelled", "timed_out", "expired", "blocked",
    },
)
if status.status == "blocked":
    # Surface a UI prompt; resume via the human-tasks API.
    ...
```

`tail_trace()` accepts the same kwarg with the same semantics.

TimeoutError handling

Both helpers raise `TimeoutError` rather than returning a partial result, because returning a "not yet terminal" status silently is easy to misuse. Pattern:

```
try:
    status = await client.platform.runs.wait_for_completion(
        run_id, timeout=300
    )
except TimeoutError:
    # Decide whether to cancel, escalate, or keep waiting with a fresh budget.
    current = await client.platform.runs.get(run_id)
    if current.status == "running":
```

```
        await client.platform.runs.cancel(run_id)
    raise
```

For `tail_trace`, events yielded *before* the timeout stay consumed:

```
seen: list = []
try:
    async for event in client.platform.runs.tail_trace(run_id, timeout=60):
        seen.append(event)
except TimeoutError:
    print(f"timed out with {len(seen)} events streamed")
    raise
```

Error handling

Source: <docs/user/errors.md>

This guide explains the typed exception hierarchy `kneo-client` raises, what each exception carries, and how to write robust catch blocks for the common operational shapes.

What "errors" mean in kneo-client

Every failure — at any layer — surfaces as a typed exception derived from `KneoError`. There is no `(ok, err)` tuple return, no `Response[T]` wrapper, no `errno` field. The standard Python `try / except` flow is the *only* error-handling shape.

Each exception carries enough context to:

- **Log the failure with traceability** — every exception has `.request_id` (the server-assigned correlation ID) and, for `POST` failures, `.idempotency_key`.
- **Branch on the operational meaning** — the exception class encodes "what went wrong" (auth vs. permission vs. server outage vs. network).
- **Read the server's reason** — `.body` is the parsed JSON the platform returned (or raw text when the response wasn't JSON).
- **Decide whether to retry, escalate, or surface** — combined with `.status` and the exception type, you can route the failure programmatically.

Hierarchy

```
KneoError
├─ KneoNetworkError          # DNS / connect / TLS / read timeout – wrapped from
httpx.HTTPError
├─ KneoBadRequestError      # HTTP 400 – semantically rejected (spec_invalid,
token_budget_exceeded, ...)
├─ KneoAuthError            # HTTP 401 – missing or invalid API key
├─ KneoPermissionError      # HTTP 403 – scope denied, or environment_policy_blocked
├─ KneoNotFoundError        # HTTP 404 – resource does not exist
├─ KneoConflictError        # HTTP 409 – state conflict, classified by .code (carries
.retry_after)
├─   └─ KneoIdempotencyMismatchError  # HTTP 409 with code idempotency_key_conflict
├─ KneoPayloadTooLarge      # HTTP 413 – body over the server's cap (carries
.max_body_bytes)
├─ KneoValidationError      # HTTP 422 – request validation failed
├─ KneoRateLimited          # HTTP 429 (carries .retry_after)
├─ KneoServerError          # HTTP 5xx – server-side failure
├─   └─ KneoServiceUnavailable  # HTTP 503 – backpressure / store down (carries .retry_after)
```

`KneoIdempotencyMismatchError` is a *subclass* of `KneoConflictError` (catching `KneoConflictError` also catches the mismatch case). Everything else is parallel.

What every exception carries

```
class KneoError(Exception):
    status: int | None          # HTTP status, or None for transport-level failures
    body: Any                  # Parsed JSON dict, raw text, or None
    request_id: str | None     # X-Request-ID echoed by the server
    idempotency_key: str | None # Idempotency-Key sent on the failing request (POSTs)
    code: str | None           # Stable snake_case code from the error envelope (kneo_serv
0.6.0+)
```

`KneoConflictError` (409), `KneoRateLimited` (429), and `KneoServiceUnavailable` (503) each add a `retry_after` field; `KneoPayloadTooLarge` (413) adds `max_body_bytes`:

```
class KneoConflictError(KneoError):
    retry_after: float | None # Seconds parsed from the Retry-After header

class KneoRateLimited(KneoError):
    retry_after: float | None # Seconds parsed from the Retry-After header

class KneoServiceUnavailable(KneoServerError):
    retry_after: float | None # Seconds parsed from the Retry-After header

class KneoPayloadTooLarge(KneoError):
    max_body_bytes: int | None # The server's configured body-size cap, when disclosed
```

`KneoServiceUnavailable` is a *subclass* of `KneoServerError` (catching `KneoServerError` also catches the 503 case). All other subclasses inherit from `KneoError` without adding fields.

Status code → exception mapping

HTTP status	Exception	When
400	<code>KneoBadRequestError</code>	Structurally valid but semantically rejected — <code>.code</code> is e.g. <code>invalid_request</code> , <code>spec_invalid</code> (diagnostics list in <code>.body</code>), or <code>token_budget_exceeded</code> .

HTTP status	Exception	When
401	<code>KneoAuthError</code>	Missing or invalid API key. Re-check the profile resolution chain.
403	<code>KneoPermissionError</code>	API key is valid but the platform won't authorize this operation. <code>.code</code> distinguishes a scope/role denial from <code>environment_policy_blocked</code> — the target environment's policy rejected the operation; fixing the key's scopes won't help there.
404	<code>KneoNotFoundError</code>	The resource (run, spec, environment, etc.) doesn't exist. Often a stale ID.
409	<code>KneoConflictError</code>	A state conflict, classified by <code>.code: run_state_conflict</code> (lifecycle fence — e.g. cancelling an already-terminal run) or <code>resource_locked</code> (held by another operation). Carries <code>.retry_after</code> . <code>idempotency_key_in_progress</code> (same key's original attempt still executing) is auto-retried by the transport per its <code>Retry-After</code> and surfaces here only once retries exhaust.
409 with code <code>idempotency_key_conflict</code>	<code>KneoIdempotencyMismatchError</code>	Same key reused with a different payload — see idempotency . Also raised for code-less legacy 409 bodies (pre-0.6.0 servers) on keyed requests.
413	<code>KneoPayloadTooLarge</code>	Request body exceeds the server's <code>KNEO_SERV_MAX_BODY_BYTES</code> cap; <code>.max_body_bytes</code> carries

HTTP status	Exception	When
		the cap when disclosed. Not retryable — shrink the payload.
422	<code>KneoValidationError</code>	Request validation failed. Either the platform's envelope (e.g. <code>.code == "guardrail_violation"</code>) or FastAPI's list-shaped detail, whose first entry is summarized into the message as <code>loc.path: msg (+N more)</code> — the full diagnostics stay on <code>.body</code> .
429	<code>KneoRateLimited</code>	Rate limit hit. <code>.retry_after</code> carries the server's hint.
503	<code>KneoServiceUnavailable</code>	Platform temporarily unavailable, classified by <code>.code: queue_full</code> (<code>kneo_serv</code> 's run-queue backpressure on <code>POST /v1/runs</code>) vs <code>store_unavailable</code> (the persistence store is down) — both carry <code>Retry-After</code> . A <code>KneoServerError</code> subclass that also carries <code>.retry_after</code> . The transport already retried within its policy honoring <code>Retry-After</code> ; one reaching your code means retries exhausted — use <code>.retry_after</code> for longer caller-side back-off.
5xx (other)	<code>KneoServerError</code>	Server-side failure. The transport already retried within its policy (for 502/503/504); a <code>KneoServerError</code> reaching your code means retries exhausted.

HTTP status	Exception	When
Other (1xx / 3xx / 4xx that aren't above)	<code>KneoError</code> (base)	Catch-all for unmodeled statuses.
Connection / DNS / TLS / read timeout	<code>KneoNetworkError</code>	Transport-level failure. The transport already retried for transient errors; a <code>KneoNetworkError</code> reaching your code means retries exhausted.

Catching patterns

Catch broadly, log richly

The most common pattern — log the full context, then decide whether to re-raise:

```
from kneo_client.core.errors import KneoError

try:
    run = await client.platform.runs.create(payload)
except KneoError as exc:
    log.error(
        "create_run failed status=%s request_id=%s idempotency_key=%s body=%r",
        exc.status,
        exc.request_id,
        exc.idempotency_key,
        exc.body,
    )
    raise
```

The `request_id` is the link to the platform's audit events — pass it along when reporting a problem to the platform operators.

Branch on specific status

When the operational meaning matters (auth flow, retry decision, user-visible error message):

```
from kneo_client.core.errors import (
    KneoAuthError,
    KneoNotFoundError,
    KneoRateLimited,
```

```

        KneoServerError,
        KneoServiceUnavailable,
    )

    try:
        run = await client.platform.runs.get(run_id)
    except KneoAuthError:
        print("API key is missing, invalid, or revoked.")
        raise
    except KneoNotFoundError:
        print(f"run {run_id!r} does not exist.")
        return None
    except KneoRateLimited as exc:
        print(f"rate-limited; server suggests waiting {exc.retry_after}s")
        await asyncio.sleep(exc.retry_after or 10)
        raise
    except KneoServiceUnavailable as exc:
        # Backpressure / overload – back off for the server's suggested window.
        # Must precede `except KneoServerError`, since it's a subclass.
        print(f"platform overloaded; backing off {exc.retry_after or 5}s")
        await asyncio.sleep(exc.retry_after or 5)
        raise
    except KneoServerError as exc:
        log.error("platform 5xx (after retries): %s", exc.body)
        raise

```

Order matters: catch more specific exceptions first (`KneoNotFoundError` before `KneoError`).

Branching on `KneoError.code`

When one status covers several operational meanings, branch on `.code` — the stable snake_case code from the platform's error envelope (`kneo_serv` 0.6.0+; `None` when the body carries no code). The code is the contract; the human-readable message is not:

```

from kneo_client.core.errors import KneoConflictError

try:
    await client.platform.runs.cancel(run_id)
except KneoConflictError as exc:
    if exc.code == "run_state_conflict":
        pass # already terminal – fine for a cancel
    else:
        raise

```

Codes worth knowing (illustrative, not exhaustive — `kneo_serv` may add new snake_case codes in a minor):

Code	Status	Meaning
<code>invalid_request</code> , <code>spec_invalid</code> , <code>token_budget_exceeded</code>	400	Semantically rejected request.
<code>environment_policy_blocked</code>	403	Environment policy rejected the operation — distinct from a scope denial.
<code>run_state_conflict</code> , <code>resource_locked</code>	409	Lifecycle fence / resource held by another operation.
<code>idempotency_key_in_progress</code>	409	Same key's original attempt still executing; auto-retried by the transport.
<code>idempotency_key_conflict</code>	409	Key replayed with a different payload → <code>KneoIdempotencyMismatchError</code> .
<code>payload_too_large</code>	413	Body over the server's cap.
<code>invalid_idempotency_key</code>	400	An <code>Idempotency-Key</code> over the server's 256-char cap (measured on the stripped value). The client validates length locally first, so this is reached only via a hand-built raw-transport request.
<code>unknown_query_parameters</code>	422	Request sent a query param the endpoint doesn't declare (<code>kneo_serv 0.11.0+</code>). The generated wrappers only send declared params, so this usually means a hand-built raw-transport call.
<code>guardrail_violation</code>	422	A guardrail blocked the run. On a <i>synchronous</i> run this is the <code>422</code> ; since <code>kneo_serv 0.11.0</code> a guardrail block also <i>terminalizes</i> the run — an

Code	Status	Meaning
		<code>async_mode=true</code> run instead ends in the <code>failed</code> terminal status (no exception; check the status).
<code>spec_path_confined</code>	422	A caller-supplied <code>spec_path</code> , <code>overlays</code> entry, or <code>skills[].source</code> resolved outside the server's allow-listed spec root (<code>KNEO_SERV_SPEC_ROOT</code>). Default-on from <code>kneo_serv</code> 1.0.0 (opt-in through <code>0.12.x</code>). The client sends these path fields unchanged — keep specs under an allow-listed root, or have the operator set <code>KNEO_SERV_SPEC_ROOT</code> . See the compatibility floor note .
<code>not_ready</code>	503	The platform isn't ready to serve yet.
<code>queue_full</code> , <code>store_unavailable</code>	503	Run-queue backpressure vs persistence-store outage; both carry <code>Retry-After</code> .

Handle idempotency-key mismatches loudly

A 409 with code `idempotency_key_conflict` means the *same key was reused with a different payload*. This is almost always a caller bug — surface it explicitly:

```
from kneo_client.core.errors import KneoIdempotencyMismatchError

try:
    await client.platform.runs.create(payload, idempotency_key=key)
except KneoIdempotencyMismatchError as exc:
    raise RuntimeError(
        f"Idempotency-Key {exc.idempotency_key!r} was reused with a different payload. "
        f"Either generate a new key for the new request or fix the payload drift."
    ) from exc
```

See [Idempotency and retries](#) for the full story on how / when this happens.

Don't catch successful-status branches

Methods on the platform / agent clients return parsed response models on success and *raise* on failure. There is no "ok / err" branching at the call site. Wrap the *call* in `try` / `except`, not the return value:

```
# Right
try:
    run = await client.platform.runs.create(payload)
    process(run)
except KneoError:
    ...

# Wrong – runs.create never returns None / False on failure; it raises
result = await client.platform.runs.create(payload)
if result is None: # never happens
    ...
```

Transport errors and network troubleshooting

`KneoNetworkError` covers everything below the HTTP layer: DNS resolution, TCP connect failures, TLS handshakes, read timeouts, connection resets. It wraps the underlying `httpx.HTTPError` as the cause — `exc.__cause__` is the original `httpx` exception if you need to inspect it.

The transport retries these automatically within `RetryPolicy.max_attempts` for transport errors and for HTTP 429 / 502 / 503 / 504. So a `KneoNetworkError` reaching your code means **all retries exhausted**:

```
from kneo_client.core.errors import KneoNetworkError

try:
    health = await client.platform.health.readyz()
except KneoNetworkError as exc:
    print(f"could not reach the platform: {exc}")
    # Treat as a hard dependency outage; don't pretend the call succeeded.
    raise
```

If you want to see the retry behavior in your logs, set the `kneo_client.transport` logger to `INFO`:

```
import logging
logging.getLogger("kneo_client.transport").setLevel(logging.INFO)
# → INFO kneo_client.transport: transport error on attempt 1; sleeping 0.20s: ...
```

Timeouts

`Profile.timeout` (default 30s, settable per-profile in `~/.config/kneo/client.toml` or via the `KNEO_TIMEOUT` env var) becomes `httpx`'s single-number timeout, which is shorthand for *all four* of `httpx`'s timeout dimensions: connect, read, write, and pool. So a `timeout=30.0` means each of those phases independently must complete within 30s, not that the whole request must finish within 30s.

A read timeout while waiting for a slow `runs.create` or `agent.specs.compile` response surfaces as `KneoNetworkError(__cause__=httpx.ReadTimeout(...))`. The transport's retry policy applies — if you're seeing these reach your code, increase `timeout` (e.g. `KNEO_TIMEOUT=120`) rather than raising `RetryPolicy.max_attempts`, since retrying a slow request just rebuilds the same wait.

For finer control (e.g. a long read budget but a short connect budget), inject a custom `httpx` client with a `httpx.Timeout` instance:

```
import httpx
from kneo_client import KneoClient

custom = httpx.AsyncClient(
    base_url=profile.url,
    timeout=httpx.Timeout(connect=5.0, read=120.0, write=30.0, pool=5.0),
)
client = KneoClient(profile=profile, http_client=custom)
# Caller owns `custom`'s lifecycle when passing http_client.
```

TLS verification

`httpx` verifies server certificates by default against the `certifi` CA bundle, and that's what `kneo-client` uses with no extra configuration. Three common operational variations:

- **Custom CA bundle** (corporate proxy with an internal CA): pass `verify="/path/to/ca.pem"` to a custom `httpx.AsyncClient`.
- **Disable verification** (development against a self-signed staging instance, *not for prod*): `verify=False`. Don't do this against an internet-reachable platform.
- **Client certificates** (mTLS-protected staging): `cert=("/path/cert.pem", "/path/key.pem")`.

All three flow through the `http_client=` injection point shown in the *Timeouts* example. There is no `kneo-client`-level wrapper for these — the configuration surface is `httpx`'s.

A failing TLS handshake surfaces as `KneoNetworkError(__cause__=httpx.ConnectError(...))` with a message like `[SSL: CERTIFICATE_VERIFY_FAILED]`. Check the `__cause__` for the precise OpenSSL error code.

HTTP proxies

httpx automatically picks up `HTTPS_PROXY`, `HTTP_PROXY`, and `NO_PROXY` from the environment — `kneo-client` inherits that behavior with no extra wiring. So:

```
HTTPS_PROXY=http://proxy.corp.example.com:8080 python my_script.py
```

works out of the box. For an explicit proxy URL set in code (without touching env vars), inject a custom httpx client:

```
custom = httpx.AsyncClient(
    base_url=profile.url,
    proxy="http://proxy.corp.example.com:8080",
)
```

A proxy that returns 502 or refuses the upstream connection surfaces as `KneoServerError` (if the proxy returns a real HTTP error) or `KneoNetworkError` (if the connection itself fails). Check `__cause__` to distinguish.

Inspecting the underlying httpx exception

For diagnosing intermittent network issues, the original httpx exception carries more context than `KneoNetworkError`'s string. Walk the cause chain:

```
try:
    await client.platform.runs.list()
except KneoNetworkError as exc:
    cause = exc.__cause__
    print(f"kneo wrapper: {exc}")
    print(f"httpx cause: {type(cause).__name__}: {cause}")
    # Common httpx exception types you'll see here:
    # httpx.ConnectTimeout    - couldn't connect within the connect budget
    # httpx.ReadTimeout       - server stopped responding mid-read
    # httpx.ConnectError      - DNS / TCP / TLS handshake failed
    # httpx.RemoteProtocolError - connection reset, partial response, etc.
    # httpx.ProxyError        - proxy-specific failure
```

When opening a support ticket for an intermittent failure, including the `__cause__`'s type and message lets the operator distinguish connection-pool exhaustion from a stuck upstream from a proxy misconfiguration.

Building error messages for users

The exceptions are designed for *internal* error handling, not for user-facing messages. If you're surfacing platform errors to end users (in a dashboard UI, a CLI prompt, etc.), build a friendly message

from the exception's attributes rather than printing `str(exc)` :

```
def user_message(exc: KneoError) -> str:
    if isinstance(exc, KneoAuthError):
        return "Your API key is invalid. Please check your credentials."
    if isinstance(exc, KneoPermissionError):
        return "You don't have permission to perform this action."
    if isinstance(exc, KneoNotFoundError):
        return "The requested item could not be found."
    if isinstance(exc, KneoRateLimited):
        wait = exc.retry_after or 60
        return f"Too many requests. Please try again in {int(wait)}s."
    if isinstance(exc, KneoServerError):
        return f"Server error (request {exc.request_id}). Please report this."
    if isinstance(exc, KneoNetworkError):
        return "Could not reach the server. Check your network connection."
    return f"Unexpected error: {exc}"
```

Why a typed hierarchy

A single `KneoError` would force callers to inspect `.status` everywhere they want to branch. A flat enum of error codes would lose the natural `isinstance` ergonomics. The hierarchy lets you catch broadly (`except KneoError`) for logging and narrowly (`except KneoAuthError`) for recovery — without losing the underlying response context, which stays attached to the exception instance.

Subclasses are added when the platform introduces a new HTTP status that warrants its own catch site, and not before. The set above is sufficient for `/v1` as it stands at `kneo_serv 1.0.0`: `0.5.0` added run-queue backpressure (`503 + Retry-After` , surfaced as `KneoServiceUnavailable`), `0.6.0` standardized the error envelope (surfaced as `.code`), and the `0.8.0` / `0.9.0` surfaces brought the `400` / `413` / `422` shapes now typed as `KneoBadRequestError` , `KneoPayloadTooLarge` , and `KneoValidationError` . `kneo_serv 0.11.0` added two new `422` codes — `unknown_query_parameters` and a terminalizing `guardrail_violation` — and `1.0.0` added a third, `spec_path_confined` (default-on spec-path confinement); all slot into the existing `KneoValidationError` via `.code` (no new subclass needed).

Reference

Helper	Where	What it does
<code>from_response(response, *, idempotency_key=None)</code>	<code>kneo_client.core.errors</code>	Maps an <code>httpx.Response</code> to the appropriate <code>KneoError</code> subclass. Used internally by

Helper	Where	What it does
		<code>Transport</code> ; rarely called directly.
<code>KneoError(message, *, status, body, request_id, idempotency_key, code)</code>	<code>kneo_client.core.errors</code>	The base. All subclasses share this constructor signature (except <code>KneoConflictError</code> , <code>KneoRateLimited</code> , and <code>KneoServiceUnavailable</code> , which add <code>retry_after</code> , and <code>KneoPayloadTooLarge</code> , which adds <code>max_body_bytes</code>).

Compatibility matrix

Source: <docs/user/compatibility.md>

This guide tells you which `kneo-client` release supports which `kneo_serv` platform version, what forward and backward compatibility mean in practice, and how to use the drop-to-transport escape hatch when you need access to an endpoint the current `kneo-client` release doesn't wrap yet.

What "compatibility" means in kneo-client

The Kneo Agent Platform's `/v1` HTTP API is a stability boundary — a `kneo-client` release pinned to one `kneo_serv` minor works against any patch-level `kneo_serv` release on the same minor line, and against newer minors that don't break `/v1`.

The pinning is *explicit* and *committed*: <schemas/openapi.json> is a `/v1`-filtered copy of one specific `kneo_serv` release's published OpenAPI spec. Bumping the pin is a deliberate PR (via scripts/bump_schemas.py), reviewed in isolation, and never happens automatically. See [ADR-004](#) for the rationale.

Current matrix

<code>kneo-client</code>	Pinned to <code>kneo_serv</code>	Tested against	Python	Status
0.9.0	v1.0.0 (<code>info.version</code> 1.0.0)	<code>kneo_serv</code> 1.0.x line	<code>>=3.12</code> (3.12 / 3.13 / 3.14 in CI)	Current
0.8.0	v0.11.0 (<code>info.version</code> 0.11.0)	<code>kneo_serv</code> 0.11.x line	<code>>=3.12</code> (3.12 / 3.13 / 3.14 in CI)	Previous
0.7.0	v0.9.0 (<code>info.version</code> 0.9.0)	<code>kneo_serv</code> 0.9.x line	<code>>=3.12</code> (3.12 / 3.13 / 3.14 in CI)	Older
0.6.0	v0.7.0 (<code>info.version</code> 0.7.0)	<code>kneo_serv</code> 0.7.x line	<code>>=3.12</code> (3.12 / 3.13 / 3.14 in CI)	Older
0.5.0	v0.5.0 (<code>info.version</code>	<code>kneo_serv</code> 0.5.x line	<code>>=3.12</code> (3.12 / 3.13 / 3.14 in CI)	Older

kneo-client	Pinned to kneo_serv	Tested against	Python	Status
	0.5.0)			
0.4.0	v0.4.0 (info.version 0.4.0)	kneo_serv 0.4.x line	>=3.12 (3.12 / 3.13 / 3.14 in CI)	Older

kneo_serv declared general availability at v1.0.0 — the /v1 HTTP API and kneo CLI are stable contracts across the 1.x line. kneo-client 0.9.0 is the GA-candidate uptake of that contract (the /v1 delta from v0.11.0 is tiny: POST /v1/specs/run gains a 202 , and a new 422 spec_path_confined code — see the floor note below).

The pinned kneo_serv version is recorded in [schemas/SOURCE.md](#) — that's the source of truth for which platform version generated the committed _generated/ tree.

How pinning works in practice

The pin is the input to the generated layer. When you bump it:

1. scripts/bump_schemas.py fetches the new openapi.json from the target kneo_serv ref (or a local checkout).
2. The spec is filtered to /v1 paths only — kneo_serv mounts every route at both /v1/... and /...; we drop the unprefixd mounts.
3. The filtered spec replaces schemas/openapi.json , schemas/SOURCE.md is updated, and _generated/ is regenerated.
4. The hand-rolled adapter layer (platform/ , agent/) may need updates if endpoints were added, renamed, or had their shapes change. The contract test [tests/contract/test_path_coverage.py](#) catches both sides of that drift.

A kneo-client minor release ships exactly one kneo_serv pin. Patches don't change pins. A pin bump can land in any minor / patch — versioning is independent.

Forward compatibility — newer kneo_serv than the pin

kneo-client ships an explicit list of every (method, path) it wraps. A newer kneo_serv that adds endpoints will still work for everything kneo-client already wraps — you just won't have wrappers for the new endpoints until the next kneo-client release.

If you need access to a new endpoint before that release, drop to the raw transport:

```
async with KneoClient.from_profile() as client:
    # Call an endpoint that doesn't have a wrapper yet:
    resp = await client._transport.request("GET", "/v1/some/new/endpoint")
    payload = resp.json()
```

The `_transport` attribute is informally accessible (single-underscore prefix). It's not part of the documented stability surface, but the API has been stable since `0.1.0` and is unlikely to churn. The transport handles auth, retries, idempotency, request-ID injection, and error mapping just like the wrapped methods do — you just lose the typed response model.

Backward compatibility — older `kneo_serv` than the pin

`kneo-client X.Y.Z` is **not** guaranteed against `kneo_serv` releases older than its pin. Wrappers may rely on response fields that older `kneo_serv` versions don't emit, and the client's error mapping assumes the current platform error shape. A `KeyError` on `from_dict()` is the typical symptom — the wrapper expects a field that wasn't in the older platform's response.

If you have to talk to an older `kneo_serv`, pin to a matching `kneo-client` minor:

Need to talk to <code>kneo_serv</code> ...	Use <code>kneo-client</code>
<code>0.12.x / 1.0.x</code>	<code>0.9.0</code> (pinned to <code>v1.0.0</code>). Takes up the GA contract: <code>async_mode</code> now honored on <code>POST /v1/specs/run (202)</code> , the <code>spec_path_confined 422</code> from default-on spec-path confinement, and consolidated client API-stability docs. A <code>v0.11.0</code> -pinned client (<code>0.8.0</code>) also works against <code>0.12.x / 1.0.x</code> for everything it wraps — <code>specs.run async</code> just surfaces the queued run via the same <code>RunCreateResponse</code> , and an out-of-root <code>spec_path</code> draws the <code>422</code> either way.
<code>0.10.x / 0.11.x</code>	<code>0.8.0</code> (pinned to <code>v0.11.0</code>). Absorbs the <code>kneo_serv 0.11.0 /v1</code> breaks (async run-create <code>202</code> , unknown-query-param <code>422</code> rejection, guardrail-block run terminalization) and adds a versioned <code>User-Agent</code> + <code>Retry-After</code> HTTP-date parsing. <code>v0.10.0</code> introduced no <code>/v1</code> change, so a <code>v0.9.0</code> -pinned client (<code>0.7.0</code>) also works against <code>0.10.x</code> for everything it wraps.

Need to talk to <code>kneo_serv ...</code>	Use <code>kneo-client</code>
<code>0.8.x / 0.9.x</code>	<code>0.7.0</code> (pinned to <code>v0.9.0</code>). Adds the skill-catalog wrapper (<code>agent.skills.list</code>), the per-request <code>skills</code> overlay on <code>runs.create</code> , <code>Page.window</code> , the typed <code>400 / 413 / 422</code> errors, and code-aware 409 classification. A <code>v0.7.0</code> -pinned client (<code>0.6.0</code>) still works against <code>0.8.x / 0.9.x</code> for everything it wraps — it just won't expose those additions.
<code>0.6.x / 0.7.x</code>	<code>0.6.0</code> (pinned to <code>v0.7.0</code>). Adds the audit-events pagination + sorting kwargs (<code>offset / sort_by / sort_order</code> , server-honored since <code>kneo_serv 0.6.0</code>), the standardized error envelope, and the <code>HumanTaskResponse.messages</code> thread projection. A <code>v0.5.0</code> -pinned client (<code>0.5.x</code>) still works against <code>0.6.x / 0.7.x</code> for everything it wraps — it just won't expose those additions.
<code>0.5.x</code>	<code>0.5.0</code> (pinned to <code>v0.5.0</code>). Because the <code>/v1</code> contract is byte-identical to <code>v0.4.0</code> , a <code>v0.4.0</code> -pinned client (<code>0.1.x – 0.4.x</code>) also works against <code>0.5.x</code> — you just won't surface the typed <code>KneoServiceUnavailable</code> for the new <code>503</code> backpressure path.
<code>0.4.x</code>	<code>0.1.x</code> , <code>0.2.x</code> , <code>0.3.x</code> , or <code>0.4.x</code> (all pinned to <code>v0.4.0</code> ; pick <code>0.4.x</code> to drop the long-deprecated <code>audit.list</code> kwargs, <code>0.3.x</code> for the additive filter kwargs on <code>runs.checkpoints / credentials.list</code> , <code>0.2.x</code> for the <code>Page + Map</code> return-type ergonomics)

(More rows added as the project ships.)

Server-version floors and behavior changes

Some wrapped surfaces depend on the server actually implementing them — the client sends the request either way, so know your deployment's `kneo_serv` version:

- **overlays are applied only from `kneo_serv 0.9.0`.** Older servers accepted the `overlays` (and `overrides / strict` on the specs routes) fields but *silently ignored* them. From `0.9.0`

they are actually applied server-side (and replayed on resume). A deployment upgrading `kneo_serv` from `0.7.x` to `0.9.x` will see previously-ignored overlay fields in stored request payloads take effect — audit those payloads before upgrading.

- **Skill catalog requires `kneo_serv >= 0.8.0`.** `client.agent.skills.list()` 404s against older servers. The per-request `skills` overlay on `runs.create` requires `>= 0.9.0`.
- **Human-task `status` filter requires `kneo_serv >= 0.8.0`.** `human_tasks.list(status=...)` is silently ignored by older servers (honored from 0.8.0, contractual in the 0.9.0 spec).
- **Spec-explain envelope requires `kneo_serv >= 0.9.0`.** The `environment / overlays / overrides` fields on `specs.explain` (including the explain leg of `dry_run`) draw a 422 from older servers.
- **Spec-path confinement is default-on from `kneo_serv >= 1.0.0`.** A caller-supplied `spec_path`, `overlays` entry, or `skills[].source` that resolves outside the server's allow-listed root (`KNEO_SERV_SPEC_ROOT`, defaulting to the server's working directory) — including absolute, `..`-traversal, `~`, or symlink-escape paths — is rejected 422 `spec_path_confined` (surfaced as `KneoValidationError` with `.code == "spec_path_confined"`). Through `0.12.x` this was opt-in (out-of-root paths only logged a deprecation warning); at `1.0.0` the warn-and-allow path is removed. The client sends these fields unchanged either way — keep your specs under an allow-listed root, or have the operator set `KNEO_SERV_SPEC_ROOT`. See [error handling](#).

Breaking changes

A breaking change in either direction triggers a major bump:

- **`kneo-client` major** — the public Python API (`kneo_client.*` namespace) changes incompatibly. Very rare.
- **`kneo_serv` major** — i.e., introduction of `/v2`. `kneo-client` may need a major bump if `/v1` is sunset; otherwise it ships a new minor that supports both.

Within a major, deprecated surfaces keep aliases for at least one minor. See the [API stability contract](#) for the full SemVer + deprecation policy.

Verifying compatibility yourself

The simplest smoke is the bundled examples. They ship in the repository (not in the wheel), so run them from a checkout:

```
# Install a specific kneo-client version
python -m pip install "kneo-client==X.Y.Z"
```

```
# Grab the MATCHING examples — check out the same release tag, or the
# examples may use surface the installed wheel doesn't have
git clone --branch vX.Y.Z https://github.com/kneo-agent/kneo-client.git && cd kneo-client

# Point it at your kneo_serv instance
export KNEO_URL=https://your-kneo-serv.example.com
export KNEO_API_KEY=...

# Run the smoke (touches health, runs, audit, agent specs, and human tasks)
python examples/01_basic_run.py YOUR_SPEC_ID
```

The seven [examples/](#) scripts collectively touch the platform health, runs, audit, agent spec, and human-task surfaces — enough to catch obvious incompatibilities in seconds.

For deeper validation, the [integration test suite](#) is env-gated; set `KNEO_TEST_URL` and `KNEO_TEST_API_KEY` and run `python -m pytest tests/integration -v`.

What the pin does *not* guarantee

The pin guarantees the *wire format* and the *path set* of `/v1`. It does not guarantee:

- **Provider availability** — whether your `kneo_serv` deployment has GPT-4 configured, an MCP server reachable, a specific runtime registered. The pin says nothing about deployment state.
- **Spec compatibility** — a spec that works on one `kneo_serv` may fail on another if the spec relies on a specific provider, tool, or platform version. Validate specs against the target environment with `client.agent.specs.validate(...)`.
- **Policy outcomes** — `policies.environment_*` queries return whatever policy is configured on the target deployment.

These are platform-deployment concerns, not client-library concerns. The client gives you a clean wire to the platform; what the platform allows is a separate dimension.